

ROSS/ROSS.Net: Rensselaer's Optimistic Simulation System User's Guide

Christopher D. Carothers – RPI CS Faculty

Mark Anderson – RPI CS PhD Student

David Bauer – MITRE

Akintayo Holder – RPI CS PhD Student

Justin Lapre – RPI CS PhD Student

Shawn Pearce – Google

Garrett Yaun – Google

Department of Computer Science

Rensselaer Polytechnic Institute

Troy, NY 12180, U.S.A.

chrisc@cs.rpi.edu

January 21, 2010

Contents

1	ROSS	4
1.1	Quick Build	4
1.2	Data Structures	5
1.2.1	Events	5
1.2.2	Logical Process Data Structure	8
1.2.3	Kernel Processes Data Structure	9
1.2.4	Processing Element Data Structure	10
1.3	ROSS Application Programming Interface (API)	11
1.3.1	Building ROSS	11
1.3.2	Main – Setting up the ROSS Model	12
1.3.3	Event Handler Data Structure	15
1.3.4	Global Variable and Function Declarations	17
1.3.5	Initialization	21
1.3.6	Event Handlers	24
1.3.7	Reverse Computation	27
1.3.8	Collecting Application Statistics/Results	29
1.4	Glossary of Functions & Variables	29
1.5	Random Number Generation Functions	31
1.6	Priority Queue	32
1.7	ROSS Errors	33
1.8	Model Development Hints	33
2	ROSS.Net – Large-scale Meta-Modeling Framework for Large-Scale Network Modeling	35
2.1	ROSS.Net API	36
2.1.1	Network Topology Definition	37
2.1.2	Network Traffic Topology Definition	38
2.1.3	Dynamic Network Link Topology Definition	38
2.2	ROSS Memory Library	38
2.3	Executing ROSS.Net	40
2.4	The Tools Directory	41
2.5	Constructing New Models	41
2.6	Model Development	42
2.6.1	LP Initialize	42
2.6.2	LP Event Handler	42
2.6.3	LP Reverse Computation Event Handler	43
2.6.4	LP Finalize	43
2.6.5	LP XML Handler	44
2.6.6	Model Main and Finalization	44
2.7	Incorporating Models into the Framework	44
2.7.1	Adding the model source code	45
2.7.2	Adding the model to the build system	46
2.8	Edit the ROSS.Net Source Code	46

2.9	ROSS.Net Modeling Paradigm	49
2.10	Model Function Hooks	49
2.11	LP Function Hooks	51
2.11.1	LP Data Types	51
2.11.2	LP XML Handler	52
2.11.3	LP Initialization	52
2.11.4	LP Event Handler	53
2.11.5	LP Reverse Computation Event Handler	54
2.11.6	LP Finalization	55
2.12	Sending Events in ROSS.Net	55

1 ROSS

ROSS is an acronym for Rensselaer’s Optimistic Simulation System. It is a parallel discrete-event simulator that executes on shared-memory multiprocessor systems. ROSS is geared for running large-scale simulation models (i.e., 100K to even 1 million object models).

The synchronization mechanism is based on Time Warp. Time Warp is an optimistic synchronization mechanism developed by Jefferson and Sowizral [10, 11] used in the parallelization of discrete-event simulation. The distributed simulator consists of a collection of *logical processes* or LPs, each modeling a distinct component of the system being modeled, e.g., a server in a queuing network. LPs communicate by exchanging timestamped event messages, e.g., denoting the arrival of a new job at that server.

The Time Warp mechanism uses a detection-and-recovery protocol to synchronize the computation. Any time an LP determines that it has processed events out of timestamp order, it “rolls back” those events, and re-executes them. For a detailed discussion of Time Warp as well as other parallel simulation protocols we refer the reader to [8].

ROSS was modeled after a Time Warp simulator called GTW or Georgia Tech Time Warp [7]. ROSS helped to demonstrate that Time Warp simulators can be run efficiently both in terms of speed and memory usage relative to a high-performance sequential simulator.

To achieve high parallel performance, ROSS uses a technique called Reverse Computation. Here, the roll back mechanism in the optimistic simulator is realized not by classic state-saving, but by literally allowing to the greatest possible extent events to be reverse. Thus, as models are developed for parallel execution, both the forward and reverse execution code must be written. Currently, both are done by hand. We are investigating automatic methods that are able to generate the reverse execution code using only the forward execution code as input. For more information on ROSS and Reverse Computation we refer the interested reader to [4, 5]. Both of these texts are provided as additional reading in the ROSS distribution.

1.1 Quick Build

Copy the following script, call it `ross-get.sh` and run it in the directory to where you would like the `rossnet` source code installed along with the `rossnet-build` directory.

```
#!/bin/sh

svn co https://subversion.cs.rpi.edu/svn/rossnet

mkdir rossnet-build
cd ross-build

# Edit ARCH to be one of x86_64 for Linux/x86 systems
#                               bgl for Blue Gene/L
#                               bgp for Blue Gene/P
#                               ppc64 for PowerPC systems
export ARCH=x86_64
```

```

# Edit CC to point to the right MPI pre-compiler script
export CC=/usr/local/mpich2-1.2/bin/mpicc

# Make sure you have Cmake install on system.
# Ross/ROSS.Net requires Cmake 2.8
cmake ../rossnet

# Build Ross models and ROSS.Net
make -j 8

# Start MPI Daemon -- delete if not required for your
# version of MPI
mpd &

# Run ROSS.Net

cd ../rossnet/trunk/rnf/kernel
ln -s ../../../../rossnet-build/trunk/rnf/kernel/rn ./rn

# Executes the TCP2 model using Optimistic synchronization
# Make sure you edit the clock rate for your processor speed.
mpirun -np 4 ./rn --scenario=tcp2 --synch=3 --clock-rate=2667000000 --nkp=512 --batch=2 --g

```

1.2 Data Structures

ROSS allows for an application programmer to design an application with the ROSS API which will simulate some real world event. The model that we use for testing and education purposes is called *pcs*, which is short for ‘personal communications services’. PCS simulates PCS/mobile phones calls being generated at a cell phone tower, and then moved around to other towers at some later time, and then the phone call ends. At the end of a phone call, another call or two may be generated. At the end of the simulation, statistics regarding how often calls were blocked (busy) due to limited tower capacity etc were collected. This details for this model can be found in [3]. This paper is included with this ROSS distribution.

ROSS utilizes three main data structures in C. There is a PE (short for processor) struct, an LP struct (short for Logical Process) and an event struct. It will be necessary to understand these structures if you are going to complete an application of your own design within ROSS.

1.2.1 Events

This is the most basic part of the ROSS system. An event is simply that, an event that occurs in the system. The ROSS engine operates on events in what we call time-stamp order. So, if we have three events with simulation timestamps of 5, 10 and 15, then the application developer would expect for the ROSS engine to process the event with

a timestamp of 5 first, then 10 second and finally 15. Successful applications will first prime the system with a few events, and then ensure that when events are processed and leave the system that new events will be created. So, with this understanding of events, we need to know what data fields within the struct are available and important to us, the application designer.

```
/*
    * tw_event:
    *
    * Holds entire event structure, one is created for each and every
    * event in use. We preallocate tw_message's into tw_event at
    * startup time so that there is no overhead of trying to locate the
    * tw_message storage area. The tw_message MUST be the last item
    * in the tw_event.
    */
```

```
DEF(struct, tw_event)
{
```

These variables are strictly for event processing and are not part of the API.

```
/*
    * next          -- Next pointer for all queues.
    * prev          -- Used by some queues.
    * up            -- Used by some queues.
    * cancel_next   -- Next event in the cancel queue for the dest_pe.
    * caused_by_me  -- Start of event list caused by this event.
    * cause_list    -- List of events caused by the event with its
    *                caused_by_me pointer set.
    */
```

```
tw_event *volatile next;
tw_event *volatile prev;
tw_event *volatile up;
tw_event *volatile cancel_next;
tw_event *volatile caused_by_me;
tw_event *volatile cause_list;
```

This is the simulation time stamp for this event. This would be 5, 10, or 15 in the example above.

```
/*
    * recv_ts       -- Actual time to be received.
    */
```

```
tw_stime      recv_ts;
```

This bit-field is used by the ROSS engine to determine in which data structure(s) the event is currently located. This is not part of the ROSS API.

```

/* Status of the event's queue location(s). */
struct
{
    BIT_GROUP(

        /* Owner of the next/prev pointers; see tw_event_owner */
        BIT_GROUP_ITEM(owner, 4)

        /* Actively on a dest_lp->pe's cancel_q */
        BIT_GROUP_ITEM(cancel_q, 1)
        BIT_GROUP_ITEM(cancel_asend, 1)

        /* Indicates union addr is in 'remote' storage */
        BIT_GROUP_ITEM(remote, 1)
        BIT_GROUP_ITEM(__pad, 1)
    )
} state;

```

The dest_lp pointer is the destination LP (logical process) for this event. We also keep track of the source LP, some state-saving information and a bit-field per event. The bit-field is available in the ROSS API for application use and ROSS does not try to impose any limits on the logical functionality of this data type. The lp_state pointer contains lp information prior to the addition of an event.

```

/*
 * dest_lp    -- Destination LP object.
 * src_lp     -- Sending LP.
 * lp_state   -- dest_lp->state BEFORE this event.
 * cv         -- Used by app during reverse computation.
 */
tw_lp        *dest_lp;
tw_lp        *src_lp;
void         *lp_state;
tw_bf        cv;

```

The tw_message pointer is critical to an event because this is where the application data is kept. ROSS provides in the API for the proper allocation of this pointer.

```

/*
 * message    -- Actual application data.
 */
tw_message   *message;

```

```
};
```

1.2.2 Logical Process Data Structure

The LP (logical Process) struct is essential to the application programmer because it provides some logical environment for events to exist in. In PCS, events are cell phones, and LPs are cell phone towers. Cell phones initiate calls which require lines on a towers. Calls move from tower to tower. When a call ends, a line in a tower is freed. If we were to model the Internet using ROSS, events might be TCP/IP packets, and LPs might be routers.

```
/*
 * tw_lp:
 *
 * Holds our state for the LP, including the lptype and a pointer
 * to the user's current state. The lptype is copied into the tw_lp
 * in order to save the extra memory load that would otherwise be
 * required (if we stored a pointer).
 *
 * Specific PE's service specific LPs, each PE has a linked list of
 * the LPs it services, this list is made through the pe_next field
 * of the tw_lp structure.
 */

DEF(struct, tw_lp)
{
```

LPs can be many different types within ROSS. The application must specify these types and that information is recorded in the *type* variable you see here. LPs also have unique ids, and are assigned to PEs (processors). The application developer must create each lp and assign it to a PE. ROSS does no dynamic load balancing, so it is left to the programmer to determine where sinks and sources should be placed. These simulations typically work best with abstracted LPs which are essentially all the same. For example, a cell phone tower has channels available for placing calls. The cell tower lp is responsible for allocating and freeing its own channels. The ROSS API provides functions for assigning and viewing this information.

```
/*
 * type      -- COPY of type from type array, holds func ptrs, etc.
 * id        -- local LP ID number
 * gid       -- global LP ID number
 * pe        -- PE that services this LP.
 */
tw_lptype    type;
tw_lpid      id;
tw_lpid      id;
tw_pe        *pe;
```


The `cur_state` pointer is used to determine the current application data for this lp. The `cur_event` pointer is the event currently being processed by this lp. These are used by the ROSS engine to determine which is the next event to process on an lp basis.

```

/*
 * cur_state      -- Current application LP data.
 * cur_event      -- Current event being processed
 * state_qh       -- Head of [free] state queue (for state saving).
 * rng            -- RNG stream array for this LP
 * type          -- Type of this LP, including service callbacks.
 */
void              *cur_state;
tw_lp_state       *state_qh;
tw_lptype         type;
tw_rng_stream     *rng;

```

The `kp` pointer is the owning kernel process for this LP. The Kernel Process data structure will be discussed next. However, from the model developers point of view, you should not be directly accessing KP data fields.

```
tw_kp             *kp;
```

This time element is the time of the last processed event in this lp. The ROSS engine uses this information to determine it's current time in the simulation and is not part of the ROSS API.

```

/*
 * last_time      -- Time of the last PROCESSED event,
 *                NOT the last SENT event
 */
tw_stime         last_time;
};

```

1.2.3 Kernel Processes Data Structure

A Kernel Process (KP) holds our state for the Kernel Process (KP), which consists only of processed event list for a collection of LPs.

```

DEF(struct, tw_kp)
{
    /*
     * id          -- ID number, otherwise its not available to the app.
     * pe          -- PE that services this KP.
     * kp_next     -- Next KP in the PE's service list.
     */

```

```

tw_kpid      id;
tw_pe        *pe;
tw_kp        *kp_next;

/*
 * last_time      -- Time of the last PROCESSED event,
 *                  NOT the last SENT event
 */
tw_stime      last_time;

```

This next series of variables is used to contain the processed event list. This is the list of events that is kept around for purposes of being able to rollback LPs. Please note, that because we are aggregating the processed event-list of many LPs, a single rollback at the KP level could result in many LPs being rolled back, so of which are not necessary. This is a trade-off in design between the need for efficient fossil collection and ‘‘false’’ rollbacks. Currently, ‘‘false’’ rollbacks do not appear to be an issue in the models we have tested to date. For more information on KPs see [5].

```

/*
 * pevent_q      -- Last event processed (for rollback).
 */
tw_event      *pevent_q;

```

And finally, we collect ROSS specific stats on a KP by KP basis.

```

/*
 * s_nevent_processed  -- Number of events processed.
 * s_e_rbs             -- Number of events rolled back by this LP.
 * s_rb_total          -- Number of total rollbacks by this LP.
 * s_rb_secondary      -- Number of secondary rollbacks by this LP.
 */
long    s_nevent_processed;
long    s_e_rbs;
long    s_rb_total;
long    s_rb_secondary;
};

```

1.2.4 Processing Element Data Structure

Most of what is contained within the PE (processor) struct is internal to the ROSS engine so it is not shown here. There is one field of note and that is the id field. Within a machine, each processor is assigned a unique PE id, starting from zero and going up. This is important to the application designer when mapping LPs to PEs.

1.3 ROSS Application Programming Interface (API)

The ROSS API is pretty small and gives the application designer a fair amount of flexibility. In this section we will discuss each of the functions available to the application programmer. We will note which functions are required where necessary. The first job of any application is to provide a main function for ROSS to be compiled against. In this main function we will need to call a couple required functions in order to setup the system properly. In this section we will use the PCS application to highlight how these functions are realized. We will start by going through the main function, and then each of the handlers (yes, this is event driven programming).

1.3.1 Building ROSS

ROSS requires the MPI library to be available on the parallel computing platform.

In the `ross` subdirectory, you will see a Makefile with the following parameters at the top:

```
# Uncomment for verbose output.
V=1

#####
# for ARCH select from:
#   1. bg1      -- Blue Gene/L using MPI
#   2. bgp      -- Blue Gene/P using MPI
#   2. x86_64   -- Linux/AMD or Intel 64 bit using MPI
#   3. ppc64    -- PowerPC 64 bit system using MPI
#
#####
ARCH=x86_64

#####
# Standard setting -- see below for definitions
#####
NETWORK=mpi
THREAD=none
GVT=mpi_allreduce
DATA_LAYOUT=byte_per_bit
QUEUE=splay
MAPPING=linear
RAND=clcg4
MEMORY=true
TIMING=true
#####
```

You only need to set the right architecture and make sure the `mpicc` pre-compiler for MPI is in the path. The other options are advanced settings and should only be changed if you are a ROSS Developer.

If you wish to increase/change the level of optimization, change the CFLAGS variable in the Makefile accordingly and shown below.

```
ifeq ($(ARCH),bgl)
    CC=mpixlc
    CFLAGS = -qflag=i:i -qattr=full -O5 -DARCH_$(ARCH)
    OPTIONS = -qtune=440 -qarch=440d
    CLOCK=blg
endif

ifeq ($(ARCH),bgp)
    CC=mpixlc
    CFLAGS = -qflag=i:i -qattr=full -O5 -DARCH_$(ARCH)
    OPTIONS = -qtune=450 -qarch=450d
    CLOCK=blg
endif

ifeq ($(ARCH),ppc64)
    CC=mpicc
    CFLAGS = -g -Wall -D_GNU_SOURCE
    CLOCK=ppc
endif

ifeq ($(ARCH),x86_64)
    CC = mpicc
    CFLAGS = -g -Wall -D_GNU_SOURCE
    CLOCK=amd64
endif
```

Just type `make clean all` and ROSS will build a library.

Next to build a sample model, such as `pcs`, change directory to `models/pcs`. Edit the CFLAGS to change the compiler optimization level and type `make`. Then to execute the `pcs` model do:

```
mpirun -np 8 ./pcs --clock-rate=2667000000.0 --batch=2 --gvt-interval=2048
```

Here, the `pcs` model will run for the default 1,000,000 time units on 8 processors which have a 2.667 GHz clock rate for cycle timer info. The batch and gvt-interval parameters determine how frequently GVT is set. See our paper [6] and others on the setting of these parameters.

1.3.2 Main – Setting up the ROSS Model

The purpose of the main function initialize the simulation model which includes the LP to PE mapping. By default it is a linear mapping. This main function builds a square mapping of LPs to PEs. It could be considered a matrix, but why make things

more complicated. Think of a map where cell phone towers are placed into a grid. Each grid space is really a processor with an lp inside of it. We may have many more cell phone towers than grid spaces, so some grid spaces will have many towers. We would like to achieve a mapping such that all grid spaces have an equal number of towers.

```
int
main(int argc, char **argv)
{
```

These are the local variables we will be using to determine the proper mapping. Because this is a grid, we will need an x and y coordinate. Also, we will have virtual x and y coordinates, xvp and yvp (for x-virtual processor and y-virtual processor). We have a simple counter, i, the number of LPs to map, TWnlp (Time Warp number of LPs), the number of processors in this machine to use, TWnpe (Time Warp number of PEs) and an lp pointer, lp (yet unallocated).

```
    int          x, y, xvp, yvp;
    int          neighbor_x[4], neighbor_y[4];
    int          vp_per_proc;
    int          i;
    int          TWnlp;
    int          TWnkp;
    int          TWnpe;
    tw_lp        *lp;
```

We have some #define variables at the top of this file to make changing the number of LPs easy to change..

```
TWnlp = NUM_CELLS_X * NUM_CELLS_Y;
```

We must have a nice machine because there appears to be four processors in it. It may really only have two processors, in which case our simulation would be hampered by context switching. The TWnpe variable should really be set to the number of processors in a machine.

```
TWnpe = 4;
```

Here we come across our first global variable, g_tw_ts_end (Global Time Warp Time Stamp Endtime). This must be a fast machine because we are telling the simulation here to run to 100 million time stamps. This is a required step if you like your simulation to end cleanly and to have statistics computed.

```
g_tw_ts_end = 100000000.0;
```

However, this can be set using the command line option:

```
--end=100000000.0
```

For a complete listing of options use --help.

This variable is the number of virtual processes per processor. The math here is not really important to us unless we are interested in doing these types of mappings. You'll notice however, that we are simply taking the number of grid spaces we desire and dividing that by the number of real processors in the system.

```
vp_per_proc = ((double)NUM_VP_X * (double)NUM_VP_Y) / (double)TWnpe;
```

We love printf's. We love them more when they are followed by a fflush(stdout);

```
printf("Running simulation for %f time units on %d processors.\n\n",
      (double) g_tw_ts_end, TWnpe);
```

This is the first function we are required to call in ROSS. The purpose of this function is to setup the ROSS internals for the simulation you desire.

```
tw_init(&argc, argv);
tw_define_lps(nlp_per_pe, sizeof(struct Msg_Data), 0);
for(i = 0; i < g_tw_nlp; i++)
    tw_lp_settype(i, &mylps[0]);
```

We pass the argc and argv parameters from main into tw_init for parsing which setups on the internals of the simulation model you desire. When then declare for each LP, what LP type it is.

```
for(x = 0; x < NUM_CELLS_X; x++)
{
    for (y = 0; y < NUM_CELLS_Y; y++)
    {
        neighbor_x[0] = ((x - 1) + NUM_CELLS_X) % NUM_CELLS_X;
        neighbor_y[0] = y;
        neighbor_x[1] = (x + 1) % NUM_CELLS_X;
        neighbor_y[1] = y;
        neighbor_x[2] = x;
        neighbor_y[2] = ((y - 1) + NUM_CELLS_Y) % NUM_CELLS_Y;
        neighbor_x[3] = x;
        neighbor_y[3] = (y + 1) % NUM_CELLS_Y;

        for (i = 0; i < 4; i++)
            Neighbors[x + (y * NUM_CELLS_X)][i] = neighbor_x[i] +
                (neighbor_y[i] * NUM_CELLS_X);
    }
}
```

After `tw_init` has parsed the commandline args, we setup our model specific mapping table. In this case, it is a simple $N \times N$ grid on which the LPs are placed. Internally within ROSS, the default mapping across processors is ‘‘linear’’. That is, we given p processors, we map an L/p slice of LPs to each processor in linear order, where L is the number of LPs for the whole model.

Now that we have the appropriate pointers to LPs and KPs, we can set it’s type. Since all PCS LP types are cell phone towers, all LP types will be `TW_CELL`, another `#define..` in this case, `TW_CELL` is defined to be 1, or, the first element in the `mylps` array. PLEASE NOTE, YOUR LP TYPE NUMBERS CANNOT START WITH ZERO. THE ZERO TYPE IS USED TO DENOTE THE END OF THE ARRAY.

Other than calling `tw_run()`, as you see below, that’s really about it to the main function. Some initializing of the simulator, and some assignment of LPs to PEs and we’re ready to go. Before we start running, we will be sure to zero out our application statistics variables.

```
/*
 * Initialize App Stats Structure
 */
TWAppStats.Call_Attempts = 0;
TWAppStats.Call_Attempts = 0;
TWAppStats.Channel_Blocks = 0;
TWAppStats.Busy_Lines = 0;
TWAppStats.Handoff_Blocks = 0;
TWAppStats.Portables_In = 0;
TWAppStats.Portables_Out = 0;
TWAppStats.Blocking_Probability = 0.0;

tw_run();
```

If we were lucky and our simulation completed without error, we can then compute and print our statistics. ROSS knows nothing about our application, so these are application defined functions. ROSS has it’s own statistics that it keeps.

```
if( tw_ismasters() )
{
    CellStatistics_Compute(&TWAppStats);
    CellStatistics_Print(&TWAppStats);
}
tw_end()

return 0;
}
```

1.3.3 Event Handler Data Structure

Next we’d like to view some of this application and try to clarify how we make something meaningful out of this simulator. We talked about LPs having types. In PCS we had

one type, TW_CELL. Here we see where that comes from.

TW_CELL is simply a #define which allows us to access the first element in the array mylps. In mylps, we have declared our event handlers. There are five major events handlers in ROSS, initialization, processing, rollback, and a final state. Remember that ROSS executes speculatively across processors, and that from time to time processors may process events out of order, in which case the system will need to be able to rollback the offending events, process the out of order event, and then begin executing again. This is why we require a rollback event handler in ROSS. We also have a state-saving handler which is not used in ROSS because state-saving is not as efficient a technique and more conservative an approach versus reverse computation.

```
#define TW_CELL 1

tw_lptype      mylps[] =
{
    {
        (init_f) Cell_StartUp,
        (event_f) Cell_EventHandler,
        (revent_f) RC_Cell_EventHandler,
        (final_f) CellStatistics_CollectStats,
        (map_f) mapping,
        sizeof(struct State)
    },
    {0},
};
```

From ross-types.h header file, the tw_lptype structure looks as follows.

```
/*
 * User implements virtual functions by giving us function pointers
 * for setting up an LP, handling an event on that LP, reversing the
 * event on the LP and cleaning up the LP for stats computation/collecting
 * results and a mapping function where take the LP gid and return
 * processor on which that LP is mapped.
 */
typedef void      (*init_f) (void *sv, tw_lp * me);
typedef int       (*event_f) (void *sv, tw_bf * cv, void *msg, tw_lp * me);
typedef void      (*revent_f) (void *sv, tw_bf * cv, void *msg, tw_lp * me);
typedef void      (*final_f) (void *sv, tw_lp * me);
typedef tw_peid   (*map_f) (tw_lpid gid);

/*
 * Each LP must have an lptype, consisting of the LP name and function
 * pointers. Each LP will contain a COPY of this data after startup,
```



```

    * allowing the LP to quickly jump into each function.
    *
    * name          -- Must be any non-zero value.
    * state_sz      -- Number of bytes that SV is for the LP.
    * init          -- LP setup routine.
    * event         -- LP event handler routine.
    * revent        -- LP RC event handler routine.
    * final         -- LP final handler routine.
    * statecp       -- LP SV copy routine.
    */
DEF(struct, tw_lptype)
{
    init_f          init;
    event_f         event;
    revent_f        revent;
    final_f         final;
    map_f           map;
    size_t          state_sz;
};

```

Alternatively, you could declare and initialize your `tw_lptype` array as follows as well in the main routine:

```

tw_lptype mylps[4];

mylps[0].init = Cell_StartUp;
mylps[0].event = Cell_EventHandler;
mylps[0].revent = RC_Cell_EventHandler;
mylps[0].final = CellStatistics_CollectStats;
mylps[0].map = mapping;
mylps[0].state_sz = sizeof(struct State);
mylps[1].init = NULL;

```

1.3.4 Global Variable and Function Declarations

Before we look at one of these handlers, lets first see what the application variables are so that we will be familiar with them when looking through the code. Below we will list every variable in `pcs` noting the intent of the application programmer.

All ROSS applications must include `ross.h` if they would like to have their application work.

```
#include <ross.h>
```

These `#defines` are application specific and determine the number of LPs in both the `x` and `y` direction. They are used, as we saw, in the mapping of LPs and KPs to PEs. In this case we see a mapping of 4 LPs (2 X 2) to 4 PEs (2 X 2).

```
#define TW_MAX_NAME_LEN 31
```

```
#define NUM_CELLS_X 2
```

```
#define NUM_CELLS_Y 2
```

```
#define NUM_VP_X 2
```

```
#define NUM_VP_Y 2
```

Channels are phone lines in a cell tower. Here we see that each cell tower will be able to support 10 simultaneous phone calls with no reserve channels.

```
#define MAX_NORMAL_CHANNELS 10
```

```
#define MAX_RESERVE_CHANNELS 0
```

We must define how large a time step to make in the simulation for certain events. Here, we define a cell phone moving from one tower to another to take 4500 time stamp units. The next call average to be 360 time stamp units, and the average call length to be 180 time step units, which in this case is equated to seconds in the real system (i.e., the average phone last 180 seconds or 3 minutes).

```
#define MOVE_CALL_MEAN 4500.0
```

```
#define NEXT_CALL_MEAN 360.0
```

```
#define CALL_TIME_MEAN 180.0
```

This is the number of initial calls to start with per cell phone tower. Since this application came from Georgia Tech, and we wanted to preserve it in ROSS, we did not try to rename variables and functions to something more meaningful.

```
#define BIG_N (double)16.0
```

Here we have more application logic. These values were picked to have logical meanings, but you can see that they are all basically integer values.

```
/*
```

```
 * When Normal_Channels == 0, then all have been used
```

```
*/
```

```
#define NORM_CH_BUSY ( !( SV->Normal_Channels & 0xffffffff ) )
```

```
/*
```

```
 * When Reserve_Channels == 0, then all have been used
```

```
*/
```

```
#define RESERVE_CH_BUSY ( !( SV->Reserve_Channels & 0xffffffff ) )
```

```
typedef int      Channel_t;
```

```
typedef int      Min_t;
```

```
typedef int      MethodName_t;
```

```

#define NONE 0
#define NORMAL_CH 1
#define RESERVE 2

#define COMPLETECALL 3
#define NEXTCALL 4
#define MOVECALL 5

#define NEXTCALL_METHOD 6
#define COMPLETIONCALL_METHOD 7
#define MOVECALLIN_METHOD 8
#define MOVECALLOUT_METHOD 9

```

This is the state of a cell tower (lp). We said earlier that our application would have to be responsible for maintaining it's own logic and we do that with this struct. This is not to be confused with state-saving. We can see here the developer is tracking the number of portables that have come through a tower, the number of attempts which are made to the tower, the number of blocks (busy signals) given, and of course the cell tower location on the grid.

```

struct State
{
    double      Const_State_1;
    int         Const_State_2;
    int         Normal_Channels;
    int         Reserve_Channels;
    int         Portables_In;
    int         Portables_Out;
    int         Call_Attempts;
    int         Channel_Blocks;
    int         Busy_Lines;
    int         Handoff_Blocks;
    int         CellLocationX;
    int         CellLocationY;
};

```

This is a reverse computation, while-one loop, struct.

```

struct RC
{
    int         w11;
};

```

Here is something we recognize from the initialization steps. We allocated space in every event we allocate for this struct, but what does it do' Here, it keeps track

of the method name to run, the time stamp at which the call moved, ended, started a new call. We also log the channel type, or phone line type and the RC struct value.

```
struct Msg_Data
{
    MethodName_t    MethodName;
    double          CompletionCallTS;
    double          NextCallTS;
    double          MoveCallTS;
    Channel_t       ChannelType;
    struct RC       RC;
};
```

Here we find a struct filled with variables for keeping tracking of statistics. When the simulation ends, we can call the statistics compute function and fill in these values for the entire system.

```
struct CellStatistics
{
    int            Call_Attempts;
    int            Channel_Blocks;
    int            Busy_Lines;
    int            Handoff_Blocks;
    int            Portables_In;
    int            Portables_Out;
    double         Blocking_Probability;
};
```

Function declarations are included to show how each function might be defined. In this document we are not going to cover each function, but only those which illustrate useful ROSS API characteristics.

```
void Cell_EventHandler(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
Min_t Cell_MinTS(struct Msg_Data *M);
void Cell_StartUp(struct State *SV, tw_lp * lp);
void Cell_NextCall(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void Cell_CompletionCall(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void Cell_MoveCallIn(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void Cell_MoveCallOut(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void RC_Cell_EventHandler(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void RC_Cell_NextCall(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void RC_Cell_CompletionCall(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void RC_Cell_MoveCallIn(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void RC_Cell_MoveCallOut(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp);
void CellStatistics_CollectStats(struct State *, tw_lp *lp);
void CellStatistics_Compute();
```

```
void CellStatistics_Print();
int Neighbors[NUM_CELLS_X * NUM_CELLS_Y][4];
```

This final variable is a global struct by which we will collect statistics at the end of the simulation.

```
struct CellStatistics TWAppStats;
```

1.3.5 Initialization

Going back to the mylp array, recall that the initialization function was called Cell_Start. When we call tw_run(), ROSS will go through each LP invoking its assigned initialization function. We are passed through a state vector and a pointer to the corresponding LP.

```
void
Cell_StartUp(struct State *SV, tw_lp * lp)
{
    int          currentcell = 0;
    int          newcell = 0;
    int          i;
    int          dest_index = 0;
    tw_stime     ts;
    struct Msg_Data TMsg;
    struct Msg_Data * TWMsg;
    tw_event *CurEvent;

    SV->Normal_Channels = MAX_NORMAL_CHANNELS;
    SV->Reserve_Channels = MAX_RESERVE_CHANNELS;
    SV->Portables_In = 0;
    SV->Portables_Out = 0;
    SV->Call_Attempts = 0;
    SV->Channel_Blocks = 0;
    SV->Handoff_Blocks = 0;
    SV->Busy_Lines = 0;
    SV->Handoff_Blocks = 0;
    SV->CellLocationX = lp->gid % NUM_CELLS_X;
    SV->CellLocationY = lp->gid / NUM_CELLS_X;
    if (SV->CellLocationX >= NUM_CELLS_X ||
        SV->CellLocationY >= NUM_CELLS_Y)
    {
        tw_error(TW_LOC, "Cell_StartUp: Bad CellLocations %d %d \n",
                SV->CellLocationX, SV->CellLocationY);
    }
}
```

This is where we prime the system with events or phone calls. GenInitPortables simply returns that BigN variable we've already seen, so we know that each LP will start out with 16 mobile subscribers/portables which are realized as events.

```
SV->Portables_In = GenInitPortables(lp);
for (i = 0; i < SV->Portables_In; i++)
{
    TMsg.CompletionCallTS = HUGE_VAL;
```

ROSS includes a reversible random number generator, and here we see the ROSS API call to generate an exponential value. By default, each LP has 1 RNG stream. By increasing `g_nRNG_per_LP`, you can change the number from a default number of RNG streams per LP. A complete listing of all random number distributions and their respective interfaces will be included at the end.

```
TMsg.MoveCallTS = tw_rand_exponential(lp->rng, MOVE_CALL_MEAN);
TMsg.NextCallTS = tw_rand_exponential(lp->rng, NEXT_CALL_MEAN);
```

We need to determine the status of this call. It may already be completed, which would be in error because we are just beginning the simulation. It may generate a new call, or move to another tower.

```
switch (Cell_MinTS(&TMsg))
{
    case COMPLETECALL:
        tw_error(TW_LOC, "APP_ERROR(StartUp): CompletionCallTS Is Min");
        break;
```

If we are generating a new call, which we need to do in order for the simulation to continue advancing, we want to create a new event and place it into the system. So we have four new, and most widely used ROSS API functions, `tw_event_new`, `tw_event_send`, `tw_event_data` and `tw_now`.

The function `tw_now` will return to us the current time (in the simulation) for an LP. Remember that this may be out of sync with other LPs because we are speculatively executing.

We create a new event by calling `tw_event_new(tw_lpid destination, tw_stime ts, tw_lp * sender)` where `sender` is the current LP and `destination` is the global ID of the destination LP. `ts` is the time stamp at which this event was sent into the *future* and is the time at which the destination LP will receive it.

We now need this new events message data pointer so that we can fill it in. We call `tw_event_data`, passing to it the event we want the `Msg_Data` pointer for, and are returned that pointer. We begin filling it in with whatever application data we like.

We then schedule the event into the future by calling `tw_event_send`. In this case, we are always sending the event to ourselves because the source and destination LPs are the same.

```

case NEXTCALL:
    ts = max(0.0, TMsg.NextCallTS - tw_now(lp));
    CurEvent = tw_event_new(lp->gid, ts, lp);
    TWMsg = (struct Msg_Data *) tw_event_data(CurEvent);
    TWMsg->CompletionCallTS = TMsg.CompletionCallTS;
    TWMsg->MoveCallTS = TMsg.MoveCallTS;
    TWMsg->NextCallTS = TMsg.NextCallTS;
    TWMsg->MethodName = NEXTCALL_METHOD;
    tw_event_send(CurEvent);
    break;

```

Here we would like to move a call to another tower, or LP. First we want to figure out where to send this event to, so we call `tw_rand_integer` (remember all LP identifiers are integers) and set that as our destination index. We then use our neighbor array to find a neighbor to this LP. Once we have done all this, we will have a global LP id for newcell. We will again create a new event with this information as the destination LP, the time the event was received into the system, and of course our source LP. It is important to note here that when determining the time stamp offset at which the new event will enter the system must be at least zero. We would not want this called to be moved backwards in time. Again we fill in the message data fields with application data and send the event to the destination LP. Because we mapped LPs to PEs, ROSS will handle the mechanism of actually moving events around in the system and processing events in time stamp order.

```

case MOVECALL:
    newcell = lp->gid;
    while (TMsg.MoveCallTS < TMsg.NextCallTS)
    {
        double result;
        currentcell = newcell;
        dest_index = tw_rand_integer(lp->rng, 0, 3);
        newcell = Neighbors[currentcell][dest_index];
        result = tw_rand_exponential(lp->rng, MOVE_CALL_MEAN);
        TMsg.MoveCallTS += result;
    }

    ts = max(0.0, TMsg.NextCallTS - tw_now(lp));
    CurEvent = tw_event_new(newcell, ts, lp);
    TWMsg = (struct Msg_Data *) tw_event_data(CurEvent);
    TWMsg->CompletionCallTS = TMsg.CompletionCallTS;
    TWMsg->MoveCallTS = TMsg.MoveCallTS;
    TWMsg->NextCallTS = TMsg.NextCallTS;
    TWMsg->MethodName = NEXTCALL_METHOD;
    tw_event_send(CurEvent);
    break;

```

```

    }
}
}

```

And that's all there is to initializing the system! At this point ROSS has 4 PEs, each with a single lp. Each lp has been primed with 16 events or more. Now it is a simple matter for ROSS to being emptying lp queues and processing events. An important attribute not to be overlooked here is that if your events to not generate at least one other event before leaving the system, the simulation may be starved and run indefinitely having not enough events to reach the end time that you specified in main. So it is important to continually generate events. Equally important is to realize that if each event generates too many events before leaving the system, then ROSS may not be able to allocate enough buffers and never progress, and the system will be considered 'choked'. So it is important for the application programmer to realize a 1:1 or 1:2 mapping of events leaving to events generated.

1.3.6 Event Handlers

The next function which is important to look at is the event handler routine. ROSS will dequeue an event from an lp struct and try to process it, or call your handler. So these handlers must work properly.

Here we see a more complex function. Its parameters include the state vector, a bit field (which is provided to your application), the message data and of course the current lp this event is scheduled onto. We simply look at the message data to determine which function we would like to call. In Cell_StartUp we set the MethodName field in all of our events, and here we exercise that information. Relatively simple. You will also note that there are a few other things going on here. First, we set the reverse-computation, while-one loop to zero. Also, we have two new API calls, tw_error and tw_exit. ROSS tries to provide some convenient way of handling errors. Every tw_error call's first parameter is TW_LOC, which is the file location of this call. Second is the string parameter, and third are the fields for the string. The tw_exit function is provided and should be used to ensure proper thread exiting on shutdown.

```

void
Cell_EventHandler(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp)
{
    *(int *)CV = (int)0;
    M->RC.wl1 = 0;

    switch (M->MethodName)
    {
    case NEXTCALL_METHOD:
        Cell_NextCall(SV, CV, M, lp);
        break;
    case COMPLETIONCALL_METHOD:

```



```

        Cell_CompletionCall(SV, CV, M, lp);
        break;
case MOVECALLIN_METHOD:
    Cell_MoveCallIn(SV, CV, M, lp);
    break;
case MOVECALLOUT_METHOD:
    Cell_MoveCallOut(SV, CV, M, lp);
    break;
default:
    tw_error(TW_LOC, "APP_ERROR(8)(%d): Invalid MethodName(%d)\n",
            lp->id, M->MethodName);
    tw_exit(1);
}
}

```

You probably thought that completing a telephone call was a simple matter, but for simulation designers it is where the work really begins. It is the point at which all the information is known.

```

void
Cell_CompletionCall(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp)
{
    int            dest_index = 0;
    int            currentcell = 0, newcell = 0;
    struct Msg_Data TMsg;
    double         result;
    tw_stime       ts;
    struct Msg_Data *TWMsg;
    tw_event       *CurEvent;

```

First we collect all our application message data.

```

    TMsg.MethodName = M->MethodName;
    TMsg.ChannelType = M->ChannelType;
    TMsg.CompletionCallTS = HUGE_VAL;
    TMsg.NextCallTS = M->NextCallTS;
    TMsg.MoveCallTS = M->MoveCallTS;

    if ((CV->c1 == (NORMAL_CH == M->ChannelType)))
        SV->Normal_Channels++;
    else if ((CV->c2 == RESERVE == M->ChannelType))
        SV->Reserve_Channels++;
    else
    {
        tw_error(TW_LOC, "APP_ERROR(2): CompletionCall: Bad ChannelType(%d) \n",

```

```

                                M->ChannelType);
        tw_exit(1);
    }
    if (SV->Normal_Channels > MAX_NORMAL_CHANNELS ||
        SV->Reserve_Channels > MAX_RESERVE_CHANNELS)
    {
        tw_error(TW_LOC, "APP_ERROR(3): Normal_Channels(%d) > MAX %d OR
                        Reserve_Channels(%d) > MAX %d \n",
                        SV->Normal_Channels, MAX_NORMAL_CHANNELS,
                        SV->Reserve_Channels,
                        MAX_RESERVE_CHANNELS);
        tw_exit(1);
    }
    TMsg.ChannelType = NONE;

```

Then we must generate new event(s) into the system.

```

switch (Cell_MinTS(&TMsg))
{
case COMPLETECALL:
    tw_error(TW_LOC, "APP_ERROR(NextCall): CompletionCallTS(%lf) Is Min \n",
            TMsg.CompletionCallTS);
    tw_exit(1);
    break;

case NEXTCALL:
    ts = max(0.0, TMsg.NextCallTS - tw_now(lp));
    CurEvent = tw_event_new(lp->gid, ts, lp);
    TWMsg = (struct Msg_Data *)tw_event_data(CurEvent);

    TWMsg->MethodName = TMsg.MethodName;
    TWMsg->ChannelType = TMsg.ChannelType;
    TWMsg->CompletionCallTS = TMsg.CompletionCallTS;
    TWMsg->NextCallTS = TMsg.NextCallTS;
    TWMsg->MoveCallTS = TMsg.MoveCallTS;

    TWMsg->MethodName = NEXTCALL_METHOD;
    tw_event_send(CurEvent);
    break;

case MOVECALL:
    newcell = lp->gid;
    while (TMsg.MoveCallTS < TMsg.NextCallTS)
    {

```

You'll notice we are incrementing the RC.wl1 variable for a move.

```

        M->RC.wl1++;
        currentcell = newcell;
        dest_index = tw_rand_integer(lp->rng, 0, 3);
        newcell = Neighbors[currentcell][dest_index];

        result = tw_rand_exponential(lp->rng, MOVE_CALL_MEAN);
        TMsg.MoveCallTS += result;
    }
    ts = max(0.0, TMsg.NextCallTS - tw_now(lp));
    CurEvent = tw_event_new(currentcell, ts, lp);
    TWMsg = (struct Msg_Data *)tw_event_data(CurEvent);
    TWMsg->MethodName = TMsg.MethodName;
    TWMsg->ChannelType = TMsg.ChannelType;
    TWMsg->CompletionCallTS = TMsg.CompletionCallTS;
    TWMsg->NextCallTS = TMsg.NextCallTS;
    TWMsg->MoveCallTS = TMsg.MoveCallTS;
    TWMsg->MethodName = NEXTCALL_METHOD;
    tw_event_send(CurEvent);
    break;
}
}

```

1.3.7 Reverse Computation

So what's left? Well, because ROSS is a speculative execution simulator, this means that we may get into trouble from time to time when LPs get ahead or fall behind one another. LPs may not always be of the same time. Some may be sinks and some may be sources. Some may never have events scheduled to them at all. ROSS guarantees the application designer that the system will run through all these inconsistencies, and without thrashing, but how does it take advantage of the rollback technology? Put simply, when ROSS realizes an out-of-order event, it runs back every event for that lp until it reaches the correct point in time when that event was to be executed. Because this out-of-order event is going to possibly change events in the future, those future events that we rolled back must be restored and rescheduled. Rescheduled is simple, but how can you restore an event without saving its original state? Simply run the formulas backwards. This is called reverse computation and is as difficult a task as it sounds, but ROSS simplifies it for the application programmer. The application programmer is required to provide ROSS with a proper rollback event handler. When ROSS rolls back an event, this handler is invoked and (hopefully) the event is restored to its original state. In order to accomplish this, ROSS provides a random number generator that is capable of being run backwards. This makes life easy for the developer, as you will see.

All we are doing here is determining which function was used when processing this event, and then calling that function's appropriate reverse computation handler.

void

```

RC_Cell_EventHandler(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp)
{
    switch (M->MethodName)
    {
        case NEXTCALL_METHOD:
            RC_Cell_NextCall(SV, CV, M, lp);
            break;
        case COMPLETIONCALL_METHOD:
            RC_Cell_CompletionCall(SV, CV, M, lp);
            break;
        case MOVECALLIN_METHOD:
            RC_Cell_MoveCallIn(SV, CV, M, lp);
            break;
        case MOVECALLOUT_METHOD:
            RC_Cell_MoveCallOut(SV, CV, M, lp);
            break;
    }
}

```

Let's look at one of those functions, specifically, the reverse computation of the function we have already seen, Cell.CompletionCall.

```

void
RC_Cell_CompletionCall(struct State *SV, tw_bf * CV, struct Msg_Data *M, tw_lp * lp)
{
    int          i;

    if (CV->c1)
        SV->Normal_Channels--;
    else if (CV->c2)
        SV->Reserve_Channels--;
}

```

Notice here that we finally use the reverse-computation, while one counter to determine how many times the random number generator was invoked. Every move call caused the RNG to be invoked twice, and so for each move call the RNG must be reversed twice. It would have been nice if this variable had been named something more meaningful, but again, we tried to maintain code integrity in order to prove results against GTW. The two applications were, for all intensive purposes, identical.

```

    for (i = 0; i < M->RC.wl1; i++)
    {
        tw_rand_reverse_unif(lp->rng);
        tw_rand_reverse_unif(lp->rng);
    }
}

```

This is all there is to it. We simply need to reverse compute the application variables we set in `Cell_CompletionCall`. Which means decrement that which was incremented. We use the lp id to rewind the random number generator so that in the future, when this event is re-executed, it will come up with the same value as the first time it was executed. Why do we care about this random number generator? It doesn't seem to be doing much for us anyway.. but it is, because you'll remember that we used the value from the random number generator to compute the destination lp of the next event.

In this way we can guarantee that not only is the simulation processing events in order through reverse computation, but also that the simulation will be deterministic. If your simulation is not deterministic, then it is worthless because it will not produce accurate, consistent results.

1.3.8 Collecting Application Statistics/Results

So at this point there is only one remaining step to complete, and it is the most important step, statistic collection. In the application PCS we define a final state handler to complete this task, then call our own functions to compute the overall statistics for the simulation run. Here we simply fill in our global struct called `TWAppStats` that we created with each LP's state information.

```
void
CellStatistics_CollectStats(struct State *SV, tw_lp * lp)
{
    TWAppStats.Call_Attempts += SV->Call_Attempts;
    TWAppStats.Channel_Blocks += SV->Channel_Blocks;
    TWAppStats.Busy_Lines += SV->Busy_Lines;
    TWAppStats.Handoff_Blocks += SV->Handoff_Blocks;
    TWAppStats.Portables_In += SV->Portables_In;
    TWAppStats.Portables_Out += SV->Portables_Out;
}
```

Then we call our application-specific statistic computation function(s) and print out the statistics as you saw at the end of the main function, after `tw_run()`. Then we `tw_exit(0)` because we are done.

1.4 Glossary of Functions & Variables

User Set Variables: The following set of variables are set by the application at start-up.

- `g_tw_ts_end` is a end time of the simulation specified as a floating point number (64 bit precision). This must be set for all (parallel or sequential) simulations. *The default value is 1.0.*
- `g_tw_mblock` is the number of events processed without going to top of the scheduler loop. This variable only has meaning when the optimistic scheduler is used. *The default value is 16..*

- `g_tw_gvt_interval` is the number of times though the main scheduler loop before computing Global Virtual Time (GVT). It combined with the previous `g_tw_mblock` control the amount of time or number of events processed between successive GVT calculations. This variable is only used when the optimistic/parallel event scheduler is used. *The default value is 16.*
- `g_tw_events_per_pe` is the number events allocated to each processor. Each processor has its own pool of event memory it manages. For sequential execution, the application must allocate enough events such that the maximum event population of the simulation model can be supported. Typically this is just equal to the number of events scheduled at start-up multiplied by the number of LPs. For parallel execution, extra memory is required over what the sequential requires, but only slightly more. We recommend `2g_tw_mblock*g_tw_mblock` as the amount extra. See [5] for a more complete discussion as to how we derived this rule of thumb. Note, this paper is included with the documentation in the ROSS distribution. *The default value is 2048.*
- `g_tw_nRNG_per_lp`: User set variable that contains the number of RNGs per LP. This is used to support multiple RNGs per LP. *The default value is 1.* See section below on random number generation.

Setup Functions: The follow functions are used to setup the model.

- `main()`: the application model must write its own main function, inside of which it initializes the simulation model.
- `tw_init(int *argc, char ***argv)`: This is the primary initialize function of the system. It requires the application to provide: number of processors, number of kernel processes (KPs), number of logical processes (LPs).
- `tw_run(void)`: Once all system initialization and LP/KP/PE mappings are established this functions transfers control to the simulation engine and completes the initialization process and then executes the simulation. The simulation run is complete when this function returns.
- `tw_end(void)`: This function computes all the ROSS specific summary statistics and outputs any model statistics as well that were not already printed from the main function.

Event Scheduling and Utility Functions: The following set of functions are used to schedule events into the future and are the core set of routines used to construct a simulation model.

- `extern tw_event *tw_event_new(tw_lpid dest, tw_stime offset_ts, tw_lp *sender)`: allocates a new event buffer and fills in the destination LP, amount of time into the future and source LP information. PLEASE NOTE, THE TIME IS AMOUNT OF TIME INTO THE FUTURE FROM TIME NOW. THUS, IT IS A RELATIVE TIME OR OFFSET TIME AND NOT AN ABSOLUTE TIME.

- `extern void *tw_event_data(tw_event * event)`: given an pointer to an event from `tw_new_event`, return the pointer to the data portion of that event buffer.
- `extern void tw_event_send(tw_event * event)`: given an pointer to an event from `tw_new_event` and its data portion is been appropriately filled in, schedule or send that event into the future at the specified time to the specified destination LP.
- `extern tw_stime tw_now(tw_lp * lp)`: given a pointer to an LP, return the current time of that LP.
- `extern void tw_error(TW_LOC, char *fmt,...)`: Generate an error and KILL the ROSS systems. The `TW_LOC` is a macro that will automatically show which SOURCE CODE file and LINE # within that file as to where the error occurred. The string denotes what the error is. PLEASE NOTE, THIS FUNCTION SHOULD NOT BE USED WHEN OPTIMISTIC/PARALLEL EXECUTION IS TURNED ON, AS ERRORS CAN OCCUR IN OPTIMISTIC PROCESSING WHICH ARE LATER ROLLED BACK. For sequential models only, it is fine to directly use this function in the simulation model code or in cases in parallel execution where you know there is an error which could not have been caused by optimistic event processing.
- `extern void tw_exit(int res)`: Terminates the simulation.

1.5 Random Number Generation Functions

The base uniform RNG is based on L'Ecuyer's Combined Linear Congruential Generator [12]. For all RNG functions, the LP identifier (`lpid`) is used to determine which set of seeds is to be used. In this simulator engine, each LP is given its own set of seeds to prevent correlations among seeds. The seeds are approximately 2^{70} calls apart. The overall period of this generator is 2^{121} . The distributions were developed based on [9].

The RNG seeds are initialized based on a seed jumping technique that is possible with this generator. The technique is extremely fast. It is much faster than reading a file of pre-generated seeds. Additionally, when you want to have 100K or even 1 million LP simulations, such seed files are not very practical.

- `g_tw_nRNG_per_lp`: User set variable that contains the number of RNGs per LP. This is used to support multiple RNGs per LP. *The default value is 1.*
- `long tw_rand_integer(tw_rng_stream * g, long low, long high)`: returns a uniform random number between the ranges of low to high inclusive. In the case of multiple RNGs per LP, you index the RNG with the LP id times the `g_tw_nRNG_per_lp` plus the RNG offset.
- `extern long tw_rand_binomial(tw_rng_stream *g, long N, double P)`: returns the number of successes (i.e., less than probability P) of the N Bernoulli trials.

- `extern double tw_rand_exponential(tw_rng_stream *g, double Lambda):` returns a exponentially distributed random number with mean, Λ .
- `extern double tw_rand_gamma(tw_rng_stream *g, double shape, double scale):` returns a Gamma distributed random number of a particular shape and scale.
- `extern long tw_rand_geometric(tw_rng_stream *g, double P):` returns the number of trials up to and including the first success (i.e., less than probability P).
- `extern double tw_rand_normal01(tw_rng_stream *g):` returns a *unit* or standard normal distributed random number, where the mean and standard deviation of the distribution are 0 and 1 respectively.
- `extern double tw_rand_normal_sd(tw_rng_stream *g, double Mu, double Sd):` returns a normal distributed random number where the mean is Mu and the standard deviation is Sd .
- `extern double tw_rand_pareto(tw_rng_stream *g, double scale, double shape):` returns a Pareto distributed random number of a particular shape. The scale parameter is used to adjust it to fit a particular mean. The mean of the distribution is $shape/shape - 1$.
- `extern long tw_rand_poisson(tw_rng_stream *g, double Lambda):` returns the number of arrivals in a given time period t assuming some average arrival rate, Λ .
- `void tw_rand_reverse_unif(tw_rng_stream *g):` this routine is used to reverse or undo the uniform part of the RNG exactly one call. Thus, for distributions that make multiple calls to the RNG, you will need to save that information and call this function that many times in order to restore the RNG seeds to the state prior to the execution of this event. This routine is to only be used for reverse execution. It should not be used in the forward execution of the simulation.

1.6 Priority Queue

ROSS supports two priority queues: Calendar Queue and Splay Tree. In a nutshell, you should use the Calendar Queue if the simulation model exhibits the following characteristics:

- Exponential, Gamma, Lognormal, or Pareto time stamp distribution.
- Few if any “ties” or simultaneous events.
- Event population remains relatively static throughout simulation run.

The simulation model fails to meet any of these characteristics, there is a strong chance the Calendar Queue will shift into running its “worst case” performance range, which is $O(n)$ for enqueue and dequeue operations. Its normal “best case” performance complexity is $O(1)$. The Splay-Tree has a rock solid $O(\log(n))$ performance complexity

for enqueue and dequeue operations and thus should be used in cases where the Calendar Queue exhibits worse-case performance.

To configure which queue is used, modify the Makefile in the *src* directory of the ROSS distribution. Here, set the *QUEUE* variable to either: *queue-calendar.o* for the Calendar Queue or *queue-splay.o* for the Splay-Tree.

For a detailed performance study on different queue algorithms, we refer the interested reader to [13].

Please note, the Heap in previous version has been deprecated. Only use the Splay-Tree as it is 2x faster than the Heap.

1.7 ROSS Errors

The following are typical errors or problems that could be generated from inside of ROSS.

- **Core Dumps:** The most common error is that your application will yield a segmentation fault/core dump. The cause of such an error usually has do with something in your application. Either initialization is not correct or insufficient event memory has been allocated. In any case, the best course of action is to put your application under a debugger (i.e., gdb) and see where the fault lies and then back track it to the source of the error. See below for more hints on model development.
- **Insufficient Event Memory.** This error only occurs when the sequential scheduler is used. Here, the simulation model has exhausted event memory. The solution is to allocate more event memory by increasing *g_tw_events_per_pe* or *g_tw_timer_reserve* if the memory usage error is timer related.
- **GVT Same for 1000 times.** This error only occurs when optimistic/ parallel execution is turned on. In this case, the simulation model has reached a state where it can no longer make progress and thus the Global Virtual Time (GVT) computation is no longer advancing. The solution here is to increase either or both the *g_tw_gvt_i* and *g_tw_mblock* parameters to allow more events to be processed per GVT epoch. Additionally, you may need to provide more optimistic memory event buffers by increasing *g_tw_events_per_pe*. If both of these fail, start looking at your simulation model.
- **Simultaneous Events.** In many discrete-event system, simultaneous events are possible even if they do not correlate to the physical system being modeled. The proper solution to handling simultaneous events is to simulate ALL possible permutations of event orderings and take an average of the outcome. As such, ROSS does not do anything special to handle tie events. That must be handled within the application. Please note, that having many tie events could result in the previously mentioned GVT advance problem.

1.8 Model Development Hints

- **Start Small:** In developing your model, keep the event population and overall size of the systems small so you can trace the event flow by hand. When you have

even 1000s of events or 100s of LPs/simulation objects, it become to much and bugs are buried to deep within the overall system.

- **Start with Sequential Execution First:** Build your model for sequential execution only first. Make sure it yields correct results across a wide variety of configurations.
- **Become Proficient with a Debugger:** The GNU debugger (GDB) or the Data Display Debugger (front end to GDB) is recommended. As with any complex software system, a debugger is an invaluable tool. However, if you are trying to find an event-flow logic error, the debugger is probably not the best method. Log files are probably better in this case, see below.
- **Construct Event Log Files:** An excellent method to debug logical errors with in a simulation model is to create LP-specific log files. This will allow complete LP-event flows to be captured. It is advisable for the developer to append the LP number to the filename.
- **Use of locks during debugging:** ROSS has a debug lock, `g_tw_debug_lck`, that can be used in the application code to provide a debugging critical section. The can be used to provide a global counter of certain kinds of activities across all processors as well as used to ensure `printf` output appears to a terminal in the correct order.
- **Forced segmentation faults:** A final technique we have found useful is to declare a pointer variable and set it to `NULL`. When an error condition is found, just dereference the pointer variable using the `*` operator and force a segmentation fault. The debugger will capture this point and allow you to start debugging at the point of the fault.

2 ROSS.Net – Large-scale Meta-Modeling Framework for Large-Scale Network Modeling

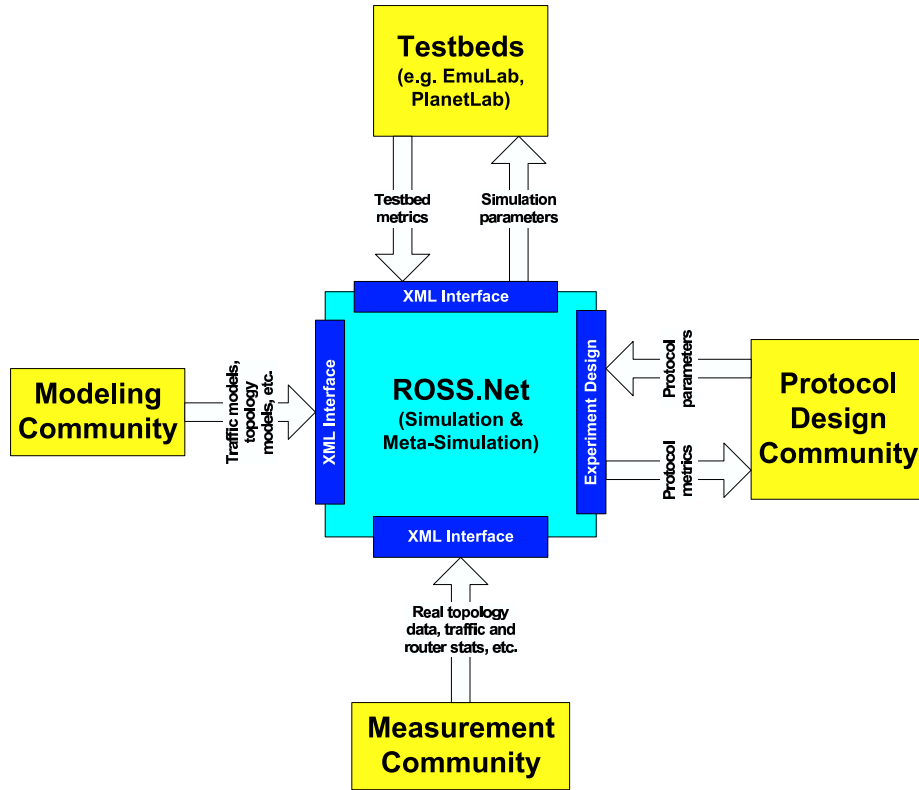


Figure 1: High level abstraction of the ROSS.Net architecture. ROSS.Net aims to bring together the four major areas of network research: simulation, protocol design, measurement and experiment design.

The ROSS.Net network simulation framework supports the study of large-scale network models on High Performance Computing (HPC) platforms. This User's Guide serves to document the features, execution and extensibility of ROSS.Net.

As shown in Figure 1, ROSS.Net aims to bring together four major areas of networking research: simulation, protocol design, network modeling and measurement and experiment design. The major components of ROSS.Net are an experiment design framework, a parallel discrete event simulator, and efficient models for network protocols and layering [14, 1, 2].

ROSS.Net incorporates the Unified Search Framework (USF) for performing optimization over the range of ROSS.Net parameters for the study of large-scale networks [15]. USF automates the process of experiment design and analysis. For simulation, at the heart of ROSS.Net is Rensselaer's Optimistic Simulation System (ROSS). ROSS is a parallel discrete event, parallel execution simulation engine that has been shown to scale to tens of thousands of processors for synthetic benchmarks, as well as realistic models of mobility and waveforms for mobile ad-hoc networks [6]. Running on top of ROSS is the ROSS.Net simulation model. The ROSS.Net simulation model provides efficient mechanisms

for layering multiple network protocol models, multiplexing packet streams, packet encapsulation, and an API interface for inter-layer communications. ROSS.Net incorporates a collection of protocol libraries (e.g.: OSPFv2, TCP-Tahoe [16], UDP, IPv4, Multicast/IPv4 etc). It allows for the rapid creation and examination of new and incremental protocol designs within an existing framework of network protocols.

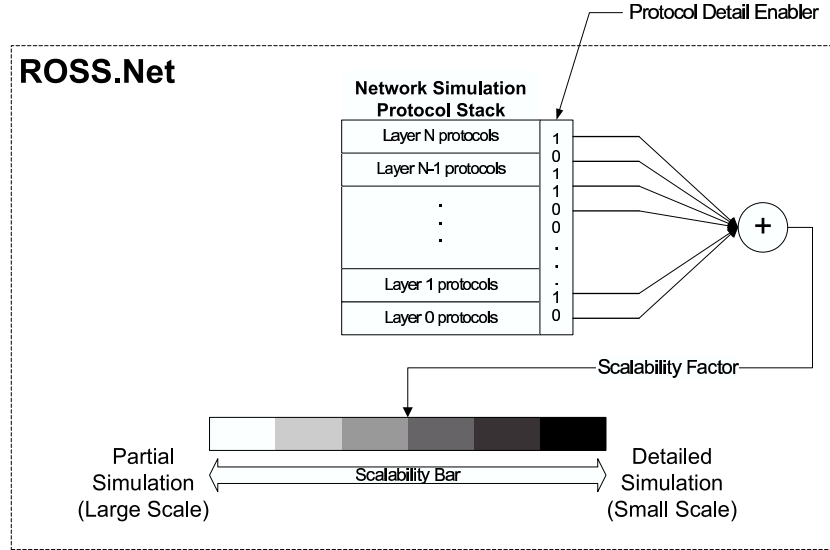


Figure 2: High level abstraction of the ROSS.Net architecture. ROSS.Net aims to bring together the four major areas of network research: simulation, protocol design, measurement and experiment design, shown in the figure on the left. The figure on the right illustrates the process of automating the design of experiments framework.

Figure 2 attempts to illustrate the modular nature of ROSS.Net, and how this framework defines the trade-off between fidelity and scalability. At the bottom of 2a, ROSS forms the basis for the framework, providing mechanisms for efficient, large-scale, parallel discrete event simulation. The ROSS.Net models include the ROSS.Net model itself as well as network protocol models. The ROSS.Net model provides a framework for wrapping and layering network protocol models. All constructs in ROSS are ROSS.Net model structures, and ROSS.Net is responsible for providing a functional API interface to the network protocol models to facilitate construction of network-specific elements in ROSS.

2.1 ROSS.Net API

ROSS.Net provides a data API in XML, defining:

- Network node / link topology
- Network traffic topology
- Dynamic link topology

Each definition is passed to ROSS.Net as a separate XML file, and ROSS.Net does some minimal error checking on the format and correctness of the files.

2.1.1 Network Topology Definition

The network topology defines networking constructs such as autonomous systems (ASes), areas, subnets and machines (or nodes). In addition, the network topology defines the initial condition of links between nodes. An XML snippet follows:

```
<rossnet>
  <as id='0' frequency='1'>
    <area id='0'>
      <subnet id='0'>
        <node id='0' links='2' type='c_router'>
          <link src='0' addr='3' bandwidth='155000000' delay='0.0015' status='up' />
          <link src='0' addr='6' bandwidth='45000000' delay='0.0015' status='up' />
          <stream port='23'>
            <layer name='ospf' level='network'>
              <interface src='0' addr='3' cost='10' />
              <interface src='0' addr='6' cost='10' />
            </layer>
          </stream>
        </node>
      </subnet>
    </area>
  </as>
</rossnet>
```

The portion of the snippet highlighted as red text is model defined XML for configuration of this element in the ospf model. ROSS.Net builds a data structure representing the network from the network topology file using the as, area, subnet, node and link XML elements. This data structure is accessible through the functional API provided by ROSS.Net:

```
// Functions to retrieve data structures
extern rn_as      *rn_getas(rn_machine * m);
extern rn_area    *rn_getarea(rn_machine * m);
extern rn_subnet  *rn_getsubnet(rn_machine * m);
extern rn_machine *rn_getmachine(tw_lpid id);
extern rn_link    *rn_getlink(rn_machine * m, tw_lpid id);

// Functions to iterate through data structures
extern rn_as      *rn_nextas(rn_as * as);
extern rn_area    *rn_nextarea(rn_area * as);
extern rn_area    *rn_nextarea_onas(rn_area * ar, rn_as * as);

extern rn_subnet  *rn_nextsn_onarea(rn_subnet * sn, rn_area * ar);
extern rn_machine *rn_nextmachine_onsn(rn_machine *, rn_subnet *);
```

One example of the use of these functions is in the OSPFv2 and IPv2 models. Each of these models computes routing tables for a given node in the network based on the topology. These protocols traverse ROSS.Net's global data structure to generate routing tables. ROSS.Net provides functions for updating and retrieving routing table entries:

```
// Retrieve route from machine m to dst
extern int rn_route(rn_machine * m, tw_lpid dst);

// Update local route from machine m to dst
extern int rn_route_change(rn_machine * m, tw_lpid dst, int new_val);
```

2.1.2 Network Traffic Topology Definition

The network traffic topology describes the traffic profile and connections between network nodes. An example would be a connection between an HTTP client (Internet browser) and server (website). The traffic topology should describe when the browser connects to website, what files are requested (corresponding to GET, PUT, POST) and when the requests are made during the simulation. Existing ROSS.Net network protocol models simply require connection definitions between clients and servers, and so the defined XML for this API is simply:

```
<connect src='0' dst='3'/>
```

Given the basic nature of the current traffic API, the connections can be specified within the `rossnet` element in the network topology definition file, or in a separate traffic definition file.

2.1.3 Dynamic Network Link Topology Definition

The dynamic network link topology describes the dynamic link topology that can occur over the simulation runtime. For example, links may go up / down at various times within the simulation:

```
<rossnet>
  <status src='0' addr='4' mode='once' change='20'/>
  <status src='15' addr='11546' mode='once' change='50'/>
</rossnet>
```

The status element indicates that the link from `src` to `addr` dynamically updates according to the function indicated by `mode`. The only mode currently supported is "once", and we use an external program to generate the dynamic link topology definition file.

2.2 ROSS Memory Library

The ROSS memory library provides a unique capability for models built in ROSS. For performance reasons, models never allocate or free memory during runtime. First and foremost, these operations would be irrevocable during reverse computation. Second,

the impact on calling these operations can make the simulation unstable and perform poorly. The preferred approach is to statically allocate all of the memory the model needs over the runtime during the initialization of the simulation.

Memory buffers are statically allocated during simulation initialization using the following function call:

```
for(i = 0; i < g_tw_nkp; i++)
    g_tcp_fd = tw_kp_memory_init(tw_getkp(i), 10000 / g_tw_nkp, sizeof(tcp_message), 1);
```

The call returns a file descriptor that is then used to signify the type of memory buffer in all future memory library calls.

The unique requirements of reverse computation caused us to generate a memory library in ROSS that facilitates reverse computation of models. The library provides functions for allocation and deallocation of memory, and reversing these operations. Notionally, the reverse computation of an allocation is to free the memory. The reverse of a free is an allocation of memory, with the data intact at the time of the free. ROSS handles all of the issues related to allocation, deallocation, and fossil collection of memory buffers at the appropriate time.

The following functions are provided by the memory buffer library in ROSS:

```
extern tw_memory    *tw_memory_alloc(tw_lp *lp, tw_fd fd);
extern void         tw_memory_free(tw_lp *lp, tw_memory *m, tw_fd fd);

extern void         tw_memory_alloc_rc(tw_lp *lp, tw_memory *head, tw_fd fd);
extern tw_memory    *tw_memory_free_rc(tw_lp *lp, tw_fd fd);

extern size_t       tw_memory_getsize(tw_kp *kp, int fd);
extern void         *tw_memory_data(tw_memory *m);

extern tw_memory    *tw_event_memory_get(tw_lp *lp);
extern void         tw_event_memory_get_rc(tw_lp *lp, tw_memory *m, tw_fd fd)

extern void         tw_event_memory_set(tw_event *e, tw_memory *m, tw_fd fd)
extern void         tw_event_memory_setfifo(tw_event *e, tw_memory *m, tw_fd fd)
```

The functions `tw_memory_alloc` and `tw_memory_free` are used during normal event processing to allocate and free memory buffers. During reverse computation, these operations can be reversed by calling the `_rc` equivalents. Notice that calling `tw_memory_free_rc` returns a pointer to the last buffer freed by the model.

The library provides two helper functions, `tw_memory_getsize`, `tw_memory_data`, that return the size of the memory buffer, and a pointer to the data portion of the memory buffer, respectively. The data portion of the memory buffer is used by the model to fill in model specific data.

Memory buffers may be attached to events, and the ROSS memory library provides functions to facilitate this feature. The functions `tw_event_memory_set` and `setfifo` attach memory buffers to events. The function `tw_event_memory_get` retrieves memory buffers from events,

and the function `tw_event_memory_get_rc` replaces buffers on events during reverse computation. Notionally, there is no action to perform in computing the inverse of setting a buffer on an event.

ROSS also provides a simple queue data structure for storing memory buffers in an LPs state. The following functions provide access to the queue data structure, and should provide the functionality to implement stacks, as well as a number of other queue-based data structures.

```
extern tw_memoryq  *tw_memoryq_init();
extern tw_memory  *tw_memoryq_pop(tw_memoryq * q);
extern tw_memory  *tw_memoryq_pop_list(tw_memoryq * q);

extern void        tw_memoryq_push(tw_memoryq * q, tw_memory * buf);
extern void        tw_memoryq_push_list(tw_memoryq *, tw_memory *, tw_memory *, int);

extern void        tw_memoryq_delete_any(tw_memoryq * q, tw_memory * buf);
extern void        tw_memoryq_fossil_collect(tw_memoryq * q, tw_lp * lp, tw_fd fd);
```

2.3 Executing ROSS.Net

ROSS.Net is executed from the command line from a UNIX shell. The command line options can be found using the `-help` option. In addition, the options for manipulating the behavior of ROSS and the individual network protocol models are given:

```
area51 ~/work/projects/rossnet: rossnet/rn --help
ROSS Network Mode: conservative
usage: rn [options] [-- [args]]
```

ROSS.Net Model:

```
--link-prob=n      link failure rate (default 1)
--bgp=n            pre-converge BGP model (default 0)
--ospf=n           pre-converge OSPFv2 model (default 0)
--log-dir=str      user specified log directory (default logs)
--topology=str     user specified XML config file (default )
--route-table=str  user specified routing table (default )
--model=str        user specified model (default )
--link-topology=str user specified link topology (default )
--scenario=str     tools scenario (default tcp)
```

ROSS MPI Kernel:

```
--read-buffer=n    network read buffer size in # of events (default 50000)
--send-buffer=n     network send buffer size in # of events (default 50000)
```

ROSS Kernel:

```
--nkp=n            number of kernel processes (KPs) per pe (default 1)
--end=ts           simulation end timestamp (default 100.00)
--batch-size=n     messages per scheduler block (default 16)
```


ROSS MPI GVT:

```
--gvt-interval=n      GVT Interval (default 16)
--report-interval=ts  percent of runtime to print GVT (default 1.00)
--help               show this message
```

2.4 The Tools Directory

The **Tools** directory contains a repository of model configurations for execution in ROSS.Net, and is located, aptly, in the `tools/` subdirectory. The primary command line option for ROSS.Net is `--scenario=value`. The value specified here indicates the sub-directory to use in the `tools/` directory where ROSS.Net will find all of the XML definition files. The network topology definition file must be named the same as the tools sub-directory, and the link topology file must be named `links.xml`:

```
area51 ~/work/projects/rossnet: ll tools/tcp/
total 20
-rw-r--r-- 1 bauerd users  22 2009-02-11 12:33 ip.rt
-rw-r--r-- 1 bauerd users  21 2009-01-22 10:06 links.xml
-rw-r--r-- 1 bauerd users 1392 2009-02-18 00:03 tcp.xml
area51 ~/work/projects/rossnet:
area51 ~/work/projects/rossnet: rossnet/rn --scenario=tcp
```

```
Opening config file: tools/tcp/tcp.xml
Opening link file: tools/tcp/links.xml
Opening routing table file: tools/tcp/ip.rt
```

```
ASes:                1
Areas:               1
Subnets:            1
Machines:            4
Links:                6
Total topology size: 848 bytes
...
```

In this example, the IPv4 protocol model reads /writes the file `ip.rt`, which contains the routing tables for the nodes, and ROSS.Net will generate output statistics for the topology, including memory consumed.

Other command line options come from ROSS, such as the simulation end time, `--end=value`.

2.5 Constructing New Models

So you want to contribute a new model to the ROSS.Net framework and test its operation in the context of the existing network protocol models? Then you need to do two things: i) develop a model and ii) hook the model into ROSS.Net framework. The following sections provide an overview on how to accomplish these two goals.

2.6 Model Development

While models written in a variety of programming languages, we focus here on C. Models must define five functions in order to operate within ROSS.Net:

1. LP initialize
2. LP event handler
3. LP reverse computation event handler
4. LP finalize
5. LP XML handler
6. Model main routine
7. Model finalization routine

The first four functions are equivalent to functions in ROSS, and handle the initialization, event processing, reverse event processing and finalization of LPs. The final three functions are specific to ROSS.Net and are hooks into the models during the XML parsing phase, and prior to and after simulation execution.

2.6.1 LP Initialize

The purpose of this function is to initialize the state associated with each LP. Per LP operations such as memory allocation, initial variable assignment, etc. are performed. A critical function of the LP initialize function is to prime the system with events. For example, an HTTP client LP may initiate a file transfer in this function. However, not all models create initial events. The IP model acts to forward IP packet events through the network, but does not have any events of its own.

2.6.2 LP Event Handler

The purpose of this function is to handle events passed into the model. Events coming into the model are specific to the model. The model must determine the processing required for one or more event types, and event types are defined by the individual models.

In addition to normal event processing, the modeler must generate code necessary for reverse computing the state of the LP. Typically this involves setting a bit in the provided bitfield for branches taken and state-saving variables into the event that contain destructive statements such as assignment. These code operations are then used later to reverse compute the state of the LP.

Beyond handling model events, a model must handle events potentially being passed between layers. Each layer model should assume that events may come from above or below the current model and handle them appropriately. The TCP model handles events being passed down from the application layer. It properly generates TCP packets that

simulate fragmentation according to the semantics of TCP. The TCP model also handles events coming from below, such as the IP layer, ensuring that data coming in from the network is in order and unfragmented. When all of the requested bytes are available, the TCP model creates and sends a final event up to the layer above it, if any exists.

Models should operate under the assumption that layers above and below exist, but they should not require knowledge from the surrounding layers. When the TCP model receives a downstream message from a layer above it, it only requires the size of the data to determine the number of TCP packets to generate. The pointer to the data is then sent with the final TCP packet. When upstream events are received from layers below TCP, the model treats these as data packets or acknowledgements. Both event types are defined by the TCP layer, and handled appropriately. When a data transmission operation is complete, the TCP model sends an upstream message to indicate completion.

ROSS.Net manages the protocol stack and determines delivery based on the direction of the event through the stack and the current model sending the event. When the top of the stack is reached, and a model sends an event up, ROSS.Net reclaims the event (no processing occurs). When the bottom of the stack is reached and a model sends an event down, ROSS.Net schedules that event to be sent between LPs in ROSS.

2.6.3 LP Reverse Computation Event Handler

The purpose of this function is to reverse compute the state of an LP during the rollback of an improperly processed event. Recall that the rollback mechanism is used to ensure that events are processed in timestamp order, and the ROSS optimistically executes events.

Generating the reverse computation code is a manual process performed by the modeler. The primary goal of this function is to restore the state of the LP prior to the current events execution. The current event is passed into this function in order to facilitate reverse computation, as is a per-event bitfield.

Reverse computation most frequently takes one of three forms. Many operations deemed constructive can be directly reverse computed. For example, incrementing a state variable during the LP event handler is matched with decrement in the reverse computation event handler. Operations deemed destructive, such as assignments or floating point operations, must be state saved into the event during normal event handling, and then the original values restored via assignment during the reverse computation. Finally, the per-event bitfield can be used to determine the path taken through the code during normal event processing.

2.6.4 LP Finalize

The purpose of this function is to provide an opportunity for each LP to process and contributed statistics for the model. Per LP statistics are often aggregated in the memory space of the model, and then the model computes some overall statistics for display at the conclusion of the simulation.

2.6.5 LP XML Handler

The purpose of this function is to configure each LP according to its corresponding XML description. ROSS.Net identifies unknown XML elements and matches them with the individual models during the setup phase of the simulation (after the LP has been initialized of course). Modelers are free to construct any schema desired, and the LPs will be given a pointer to their node in the XML document.

2.6.6 Model Main and Finalization

ROSS.Net contains the main function, and so these two functions are used to give each layer model the opportunity to perform some processing before and after the execution of the simulation. In particular, the model main functions are called after ROSS internal data structures have been configured, but before the LP initialization and XML handling phases. The purpose is to give the model a chance to initialize global memory prior to these operations for use by the LPs. The model finalization functions is provided to give each model the opportunity to report any final results, and is called immediately after the conclusion of the simulation.

2.7 Incorporating Models into the Framework

NOTE: Will be depreciated in the next release – we are moving to a CMAKE build framework.

The build system for ROSS.Net was originally chosen because of a lack of support for shared libraries on some supercomputing systems (e.g., Cray XT4) and because static libraries with unused code can impact the performance of an application. The autoconf and libtool build tools provide a great deal of facilities that allow for the dynamic incorporation of models at compile time into the creation of a minimal executable. Unfortunately, the high learning curve associated with these tools means that the build system for ROSS.Net is not yet complete. In particular, models must still be manually inserted into the ROSS.Net code.

The autoconf tool allows for configurable options to be specified during the configuration phase of the software, including creating preprocessor definitions in the code.

The steps are:

1. add the model source code to the system
2. add the model to the build system
3. edit the ROSS.Net code
4. (re-)make the software

In the following sections we will outline these steps in detail for a new model named mymodel.

2.7.1 Adding the model source code

The source code for all models is contained in the modules subdirectory. Within this directory are a series of subdirectories indicating a notional layering of protocol models. Choose the most appropriate place and name for your source code.

```
area51 ~/work/projects/rossnet: ll modules/
total 68
drwxr-xr-x 14 bauerd users 4096 2009-02-12 15:30 application
drwxr-xr-x  6 bauerd users 4096 2009-02-12 15:30 datalink
drwxr-xr-x 11 bauerd users 4096 2009-02-12 15:30 network
drwxr-xr-x  5 bauerd users 4096 2009-02-12 15:30 physical
drwxr-xr-x  2 bauerd users 4096 2009-02-12 15:30 presentation
drwxr-xr-x  2 bauerd users 4096 2009-02-12 15:30 session
drwxr-xr-x  5 bauerd users 4096 2009-02-12 15:30 transport
area51 ~/work/projects/rossnet: ll modules/application/
total 84
drwxr-xr-x 2 bauerd users 4096 2009-02-06 13:18 mymodel
drwxr-xr-x 2 bauerd users 4096 2009-02-06 13:18 bootp
drwxr-xr-x 2 bauerd users 4096 2009-02-06 13:18 dns
drwxr-xr-x 4 bauerd users 4096 2009-02-18 13:01 ftp
drwxr-xr-x 2 bauerd users 4096 2009-02-06 13:18 http
drwxr-xr-x 3 bauerd users 4096 2009-02-12 15:30 ip-traffic
drwxr-xr-x 2 bauerd users 4096 2009-02-06 13:18 rlogin
drwxr-xr-x 2 bauerd users 4096 2009-02-06 13:18 smtp
drwxr-xr-x 2 bauerd users 4096 2009-02-06 13:18 snmp
```

ROSS.Net does not use in any way the directory structure to determine the layering of models; layering follows the description in the XML topology definitions.

The final step is to create the file Makefile.am in your subdirectory indicating your model source and header files, and to edit the Makefile.am in the directory above your model:

```
area51 ~/work/projects/rossnet: cat modules/application/Makefile.am
SUBDIRS=\
    ftp \
    mymodel

area51 ~/work/projects/rossnet: cat modules/application/mymodel/Makefile.am
lib_LTLIBRARIES=libmymodel.la

AM_CFLAGS=-g -Wall

include_HEADERS=\
    mymodel-extern.h \
    mymodel-types.h \
    mymodel.h
```

```
libmymodel_la_SOURCES=\
    mymodel-global.c \
    mymodel-xml.c \
    mymodel.c
```

2.7.2 Adding the model to the build system

The first step in adding a new model to the current build system is to edit the file `configure.in`, located at the root directory. The following lines must be added:

```
AC_ARG_WITH(mymodel,
    [ --with-mymodel          Use mymodel module],
    AC_DEFINE(HAVE_MYMODEL_H, 1,    mymodel module))
```

And at the bottom of the file, the location of the model must be specified so that the Makefile can be constructed for your model:

```
AC_CONFIG_FILES([Makefile
    rossnet/Makefile \
    modules/Makefile \
    ...
    modules/network/ip/Makefile \
    modules/transport/Makefile \
    modules/application/mymodel/Makefile \
    modules/transport/tcp/Makefile])
```

Then simply run the bootstrap script to regenerate the configure program. When running `configure`, the command line option `--with-mymodel`. While all of the source code will be compiled, each model is compiled into a static library, and the final executable is built from linking only those libraries specified in the configuration.

2.8 Edit the ROSS.Net Source Code

The ROSS.Net source code must be manually updated to include the hooks for each model added, since we do not support shared libraries.

The list of files to be modified include:

- `rn_config.h`
- `rn.c`

Editing the file `rn_config.h`, the header files for the new model must be included. In order to encapsulate and segregate the new model from the ROSS.Net source code, the recommendation is to create a single header file for the new model that includes all other header files for the model. Then as new header files are added to the model, only the main model header file requires updating.

```

#ifndef INC_rn_config_h
#define INC_rn_config_h

#ifdef HAVE_FTP_H
#include "../modules/application/ftp/ftp.h"
#endif

#ifdef HAVE_TCP_H
#include "../modules/transport/tcp/tcp.h"
#endif

#ifdef HAVE_MMODEL_H
#include "../modules/application/mymodel/mymodel.h"
#endif

...

#endif

```

Then the contents of mymodel.h may have the following content:

```

#ifndef INC_mymodel_h
#define INC_mymodel_h

#include <rossnet.h>

#include "mymodel-types.h"
#include "mymodel-extern.h"

#endif

```

Editing the file rn.c involves informing the ROSS.Net source code what the names of the function hooks are, and incrementing a single variable in ROSS. The object in rn.c that must be modified is:

```

rn_lptype      layer_types[] = {

#ifdef HAVE_OSPF_H
    {
        (init_f) ospf_startup,
        (event_f) ospf_event_handler,
        (revent_f) ospf_rc_event_handler,
        (final_f) ospf_statistics_collection,
        (map_f) NULL,
        sizeof(ospf_state),
        "ospf",
        (rn_xml_init_f) ospf_xml,
        (md_init_f) ospf_main,

```

```

        (md_final_f) ospf_md_final}
    ,
#endif
#ifdef HAVE_MYMODEL_H
    {
        (init_f) mymodel_init,
        (event_f) mymodel_event_handler,
        (revent_f) mymodel_rc_event_handler,
        (final_f) mymodel_final,
        (map_f) NULL,
        sizeof(mymodel_state),
        "epi",
        (rn_xml_init_f) mymodel_xml,
        (md_init_f) mymodel_main,
        (md_final_f) mymodel_md_final}
    ,
#endif
...
    {0}
    ,
};

```

The second modification to rn.c occurs in the main function:

```

int
main(int argc, char **argv, char **env)
{
    ...
#ifdef HAVE_TCP_H
    g_tw_memory_nqueues += 1;
#endif

#ifdef HAVE OSPF_H
    g_tw_memory_nqueues += 5;
#endif

#ifdef HAVE_MYMODEL_H
    g_tw_memory_nqueues += 1;
#endif
    ...
    return 0;
}

```

Your model will likely contain at least one memory buffer type, if not more. The TCP model defines one data structure for all event types, using one memory buffer queue. However, the OSPF model defines multiple event types that vary greatly in size, and so allocates five memory queues, one for each event type, to ensure that the memory allocated is optimal to the model.

Memory buffers will be covered in more detail in Section 2.9.

At this point, assuming the source code files exist and has been configured, the system can be built by typing make.

2.9 ROSS.Net Modeling Paradigm

The ROSS.Net paradigm was designed to match closely with the modeling environment provided by the ROSS simulator. Models built in ROSS and ROSS.Net follow the LP paradigm, where LPs model processes, and events are passed between LPs model actions occurring in the system. A simple example is a queuing system, where LPs model service queues, and events passing between queues convey information for processing at the queue. The IP layer in ROSS.Net implements a simple queuing model for links between nodes, and IP packets are sent between IP LPs. When an IP packet arrives at an IP LP, that LP then determines which link to use to forward the packet, and forwards the packet to the LP at the other end of the link. The timestamp of the forwarded packet is chosen based on the link bandwidth and delay, and the current load on the link. Queuing occurs when link bandwidth is consumed. Then the forwarded packet occurs at a time equal to the amount of time for the packet to traverse the link, plus the time of the last sent packet. The time of the last sent packet on the link is then updated for the current packet.

To build a model in ROSS.Net, the functions from Section 2.6 must be defined. In this section we will illustrate the process of developing a model in ROSS.Net by illustrating the TCP model in ROSS.Net.

2.10 Model Function Hooks

From Section 2.6, the model function hooks are:

1. Model main routine
2. Model finalization routine

The TCP source code defines these hooks in tcp-model.c. During the call to the model main function we are in the initialization phase of the simulation, and ROSS.Net has processed the XML definitions, generated global data structures for the topology and initialized the ROSS data structures. However, the initialization of LPs has yet to take place. This is the point at which models have an opportunity to read configuration files, create global variables, add command line options, and initialize ROSS memory buffer queues.

```
const tw_optdef tcp_opts[] =
{
    TWOPT_GROUP("TCP Tahoe Model"),
    TWOPT_UINT("tcp-log", g_tcp_log_on, "turn on logging"),
    TWOPT_END()
};
```

```

void
tcp_md_init(int argc, char ** argv, char ** env)
{
    int      i;

    tw_opt_add(tcp_opts);

    g_tcp_stats = tw_calloc(TW_LOC, "", sizeof(*g_tcp_stats), 1);

    if(tw_ismaster() && g_tcp_log_on)
    {
        g_tcp_f = fopen("tcp.data", "w");

        if(!g_tcp_f)
            tw_error(TW_LOC, "Unable to open TCP validation file!");
    }

    for(i = 0; i < g_tw_nkp; i++)
        g_tcp_fd = tw_kp_memory_init(tw_getkp(i), 1000000 / g_tw_nkp,
                                      sizeof(tcp_message), 1);

    if(tw_ismaster())
    {
        printf("\nInitializing Model: TCP Tahoe\n");
        printf("\t%-50s %11d (%1d)\n",
              "TCP Mbufs", 1000000, g_tcp_fd);
    }
}

void
tcp_md_final()
{
    if(!tw_ismaster())
        return;

    fclose(g_tcp_f);

    printf("\nTCP Model Statistics: \n\n");
    printf("\t%-50s %11d\n", "Sent Packets", g_tcp_stats->sent);
    printf("\t%-50s %11d\n", "Recv Packets", g_tcp_stats->recv);
    ...
}

```

The first thing we see is the definition of the TCP model command line options, `tcp_opts`. ROSS provides a simple mechanism for models to add their own command line arguments, and ROSS will process these arguments properly, setting the variable to the user supplied value, or keeping the default value if no user value defined, and printing the values

as part of the `--help` option.

In `tcp_md_init` the model performs the initialization of global variables, and passes the command line arguments for the model to ROSS via `tw_opt_add(..)`. Most importantly, the model defines the size and quantity of memory buffers required by the model.

During the call to the model finalization function, `tcp_md_final`, the simulation has completed, and the LP finalization function has been called for all LPs. This function gives the model a final opportunity to output results and statistics, and perform any final logging.

ROSS provides a helper function, `tw_ismaster()`, that returns true for a single CPU in the processor configuration. This function can be used for the gather phase of the model, serializing data collection and output.

2.11 LP Function Hooks

The TCP model defines the four primary LP function hooks in `tcp.c`, and the per LP XML handler function in `tcp-xml.c`. The four primary hooks are the heart of the model, and define the process that will occur at each stage of the LP: init, event processing, reverse computation event processing, and finalization.

2.11.1 LP Data Types

The data types for the TCP model are defined in the file `tcp-types.h`. The two primary types define the state for TCP LPs, and TCP memory buffers:

```
DEF(struct, tcp_state)
{
    int            unack;
    int            seq_num;
    int            dup_count;
    int            len;
    int            timeout;
    int            rtt_seq;
    short          out_of_order[33];
    int            rto_seq;
    int            mss;
    int            recv_wnd;
    ...
    tcp_statistics *stats;
};

DEF(struct, tcp_message)
{
    int            ack;
    int            seq_num;
    tw_lpid        src;
};
```

The model defines the LP state variables required for processing events, and the structure of data for TCP events.

2.11.2 LP XML Handler

The first per LP function called is the LP XML handler. This allows each LP to receive its configuration from the XML definitions prior to their initialization functions being called.

Each model defines and handles its own XML elements per LP type:

```
<layer name='tcp' level='transport'>
  <mss>4096</mss>
  <recv_wnd>32</recv_wnd>
</layer>
```

Then the model defines an XML handler that ROSS.Net will call during the initialization phase. ROSS.Net provides pointers to the state of the LP, the XML node pointer, and the LP object.

```
void
tcp_xml(tcp_state * state, const xmlNodePtr node, tw_lp * lp)
{
    xmlNodePtr      e;

    state->mss = g_tcp_mss;
    state->recv_wnd = g_tcp_rwd;

    for(e = node->children; e; e = e->next)
    {
        if(0 == strcmp((char *) e->name, "mss"))
            state->mss = atoi((char *)e->children->content);

        if(0 == strcmp((char *) e->name, "recv_wnd"))
            state->recv_wnd = atoi((char *)e->children->content);
    }
}
```

The behavior of the model is to use the XML element to define each LPs state variables, or to use the global variables if not specified.

2.11.3 LP Initialization

The second function call is the LP initialization function. This function defines the implementation of the per LP initialization:

```
void
tcp_init(tcp_state * state, tw_lp * lp)
{
```

```

state->cwnd = 1;
state->ssthresh = TCP_SND_WND;
state->rto = 3.0;
state->connection = rn_getmachine(lp->gid)->conn;

state->stats = tw_calloc(TW_LOC, "", sizeof(tcp_statistics), 1);

...
if(-1 != state->connection)
{
    rn_message m;
    m.size = 40960;
    start_transfer(state, NULL, &m, lp);
}
}

```

The LP state variables are allocated and initialized and per LP data structures created. The function `rn_getmachine(..)` is called to retrieve the id of the destination LP connected to this LP. Rather than relying on some application model to pass events down into the TCP layer, this TCP model is designed to start a file transfer immediately upon initialization.

2.11.4 LP Event Handler

The TCP LP event handler function defines the computation that occurs when events are received at TCP clients / servers. The function is defined by the code:

```

void
tcp_event_handler(tcp_state * state, tw_bf * bf, rn_message * m, tw_lp * lp)
{
    tw_memory      *b;
    tcp_message     *in = NULL;

    // if no membuf, then must be application layer
    if(NULL != (b = tw_event_memory_get(lp)))
        in = tw_memory_data(b);

    if(!b)
        tw_error(TW_LOC, "no membuf on TCP event!");

    if(m->type == DOWNSTREAM)
        if(m->size)
            start_transfer(state, bf, m, lp);
        else
            stop_transfer(state, bf, m, lp);
    else if(m->type == TIMER)
        tcp_timeout(state, bf, in, lp);
    else if(state->connection != -1)
        tcp_process_ack(state, bf, in, lp);
}

```

```

        else
            tcp_process_data(state, bf, in, lp);

        if(b)
            tw_memory_free(lp, b, g_tcp_fd);
    }

```

The ROSS.Net event handler function prototype provides pointers to our current LPs state, the current events bitfield (for reverse computation), the current event, and our own LP object. The objects are used in subsequent calls to ROSS and ROSS.Net functions.

The first step in the model is to retrieve our TCP event data from the memory subsystem, via the calls to `tw_event_memory_get(..)` and `tw_memory_data(..)`. This function retrieves a memory buffer previously allocated and attached to the current event by the TCP model.

Inspecting the type ROSS.Net event received reveals the processing required. Primarily, the ROSS.Net event type is one of `TIMER`, `UPSTREAM` or `DOWNSTREAM`. If the event is of type `DOWNSTREAM`, then we can assume there is a model layer above ours, and the TCP model starts a new file transfer based on the size of the event passed in. Events of type `TIMER` can be constructed by protocol models to implement operating system alarms. Actual TCP implementations set an alarm called Receiver Timed Out (RTO). The idea is that when TCP sends data to a client, and that client fails to respond with an acknowledgement (ACK) within a set amount of time, then the TCP server should attempt to resend the data packet.

Finally, if none of the above cases are true, then we are either a TCP client receiving data, or a TCP server receiving an ACK. This can be determined based on the value of the LP state variable `connection`, configured in the TCP LP init function. The TCP model defines processing for each event type in other source files.

At the completion of event processing, the final step is to free the memory buffer associated with the event. This is an important step, because it allows the memory buffer to be accounted for by the simulation system. It will become available only once ROSS determines that the current event and memory buffers are safe to reclaim.

2.11.5 LP Reverse Computation Event Handler

The TCP LP reverse computation event handler function defines the reverse computation that occurs when events are being rolled back by the simulator. The primary purpose of this function is to give the model a chance to restore the state of the LP prior to processing the current event. The function is defined by the code:

```

void
tcp_rc_event_handler(tcp_state * state, tw_bf * bf, rn_message * m, tw_lp * lp)
{
    tw_memory      *b;
    tcp_message     *in;

    if(NULL == (b = tw_memory_free_rc(lp, g_tcp_fd)))
        tw_error(TW_LOC, "No membuf on event!");
}

```

```

    in = tw_memory_data(b);

    if(m->type == TIMER)
        tcp_rc_timeout(state, bf, in, lp);
    else if(state->connection != -1)
        tcp_rc_process_ack(state, bf, in, lp);
    else
        tcp_rc_process_data(state, bf, in, lp);

    tw_event_memory_get_rc(lp, b, g_tcp_fd);
}

```

As with the event handler, we must first determine the type of the current event so that we can compute the reverse code and restore the LP state. Following the same logic, the event must be either the RTO timer, data from a TCP server, or an ACK from a TCP client. If the event being rolled back is from the layer above us, then no action is required as we return to the ‘‘start’’ state of the LP.

Memory buffers must be reverse computed as well, and the ROSS memory buffer library provides API calls to simplify this process. First, we must un-free the event. At the conclusion of the reverse event handler, we must re-attach the memory buffer to the event so that when forward processing resumes, the memory buffer is available.

2.11.6 LP Finalization

The TCP LP finalization functions gathers per LP statistics into a global variable and logs output (not shown):

```

void
tcp_final(tcp_state * state, tw_lp * lp)
{
    ...
    g_tcp_stats->sent += state->stats->sent;
    g_tcp_stats->recv += state->stats->recv;
    g_tcp_stats->ack_sent += state->stats->ack_sent;
    g_tcp_stats->ack_recv += state->stats->ack_recv;
    g_tcp_stats->ack_invalid += state->stats->ack_invalid;
    g_tcp_stats->bad_msgs += state->stats->bad_msgs;
    g_tcp_stats->dropped_packets += state->stats->dropped_packets;
    ...
}

```

2.12 Sending Events in ROSS.Net

Sending events in ROSS.Net is slightly different than in ROSS because ROSS.Net provides functionality required by most if not all network protocol models. For purposes of demonstration, the TCP model defines a helper function for sending events in the TCP model:

```

void
tcp_event_send(tcp_state * state, tw_lpid src, tw_stime ts, tw_lpid dst, int size, int seq_num, int
{
    tw_event      *e;
    tw_memory      *b;

    tcp_message     *m;

    e = rn_event_new(dst, ts, lp, DOWNSTREAM, 8.0 * size);
    b = tw_memory_alloc(lp, g_tcp_fd);
    m = tw_memory_data(b);

    m->src = src;
    m->ack = ack;
    m->seq_num = seq_num;

    tw_event_memory_set(e, b, g_tcp_fd);

    rn_event_send(e);
}

```

First, the TCP model calls `rn_event_new`, supplying the destination LP's id, the offset timestamp, a pointer to the current LP, the direction of the event in the protocol stack, and the size of the data to be sent.

The destination LP is determined by the LPs connection from the Traffic topology. If there are no layers defined in the XML description beneath TCP LPs, then the event will be sent directly to the destination LP. If models such as the IP model are defined below TCP, the TCP packets will be routed through the IP router network. The TCP model need not be aware of model layers below TCP, ROSS.Net will ensure proper delivery of the event according to the XML definitions.

Whether the TCP model is sending an acknowledgement, or data packet, the offset timestamp is always 0.0. This TCP implementation does not account for queuing delays on the packet, instead relying on the layering of protocol models to provide queuing semantics based on the size of the message sent.

Whether an acknowledgment or data, every TCP packet flows down through the protocol stack, or DOWNSTREAM. When a TCP client receives a complete event, this is the only time the TCP model does not use this function. Instead, the TCP model defines code specific to that purpose elsewhere in the code. When sending an UPSTREAM event, the code is simply:

```

rn_event_send(rn_event_new(lp->gid, 0.0, lp, UPSTREAM, state->len));

```

Here the source and destination are the same, the timestamp offset is 0.0, the direction is UPSTREAM, and the size of the data being sent is the full size of the file received over the course of the transmission. No memory buffer is attached to the event, as any layers above the TCP layer have no knowledge of the TCP memory buffer structure. If the application on the sending side provided an application memory buffer, then

the TCP model should provide that memory buffer on the receivers side when the send completes. This is not currently implemented in this TCP model. This functionality could be implemented by simply having the TCP server attach the application memory buffer on an event and sending it to the TCP client. The TCP client would then forward the buffer up the stack to the application.

The TCP model allocates a memory buffer from the ROSS system, fills in the data portion of the packet, and attaches the memory buffer to the event by calling `tw_event_memory_set(...)`. Each protocol model defines and uses its own memory buffers for application data, and attaches them to newly created events, and pulls them off events during event processing.

Finally, the TCP model sends the event. The event will be passed down to lower protocol models, if any are defined on this node, and upon reaching the bottom of the stack the event is injected into ROSS.

References

- [1] D. W. Bauer, G. Yaun, C. D. Carothers, M. Yuksel, and S. Kalyanaraman, ‘‘A case study in meta-simulation design and performance analysis for large-scale networks,’’ in *Proceedings of the 2003 Winter Simulation Conference*, 2003.
- [2] D. Bauer, G. Yaun, C. Carothers, M. Yuksel, and S. Kalyanaraman, ‘‘A case study in meta-simulation design and performance analysis for large-scale networks,’’ in *Proceedings of the 36th conference on Winter simulation*. Winter Simulation Conference, 2004, pp. 206--214.
- [3] C. D. Carothers, R. M. Fujimoto and Y-B. Lin, A case study in simulating PCS networks using time warp. In *Proceedings of the 9th Workshop on Parallel and Distributed Simulation*, June 1995, pages 87--94.
- [4] C. D. Carothers, K. Perumalla and R. M. Fujimoto. Efficient Parallel Simulation Using Reverse Computation *ACM Transactions on Modeling and Computer Simulation*, volume 9, number 3, July 1999.
- [5] C. D. Carothers, D. Bauer and S. Pearce. High-Performance, Low Memory, Modular Time Warp System, In *Proceedings of the 14th Workshop of Parallel and Distributed Simulation (PADS 2000)*, pages 53--60, May 2000.
- [6] C. D. Carothers, D. W. Bauer, and S. O. Pearce, ‘‘Ross: A high-performance, low memory, modular time warp system,’’ *Journal of Parallel and Distributed Computing*, 2002.
- [7] S. Das, R. Fujimoto, K. Panesar, D. Allison, and M. Hybinette. GTW: A Time Warp system for shared memory multiprocessors. In *1994 Winter Simulation Conference Proceedings*, pages 1332--1339, December 1994.
- [8] R. M. Fujimoto. Parallel discrete-event simulation. *Communications of the ACM*, 33(10):30--53, October 1990.

- [9] R. Jain. *The Art of Computer Systems Performance Analysis*, John Wiley and Sons, Inc. New York, 1991.
- [10] D. R. Jefferson and H. Sowizral. Fast concurrent simulation using the Time Warp mechanism, part I: Local control. Technical Report N-1906-AF, RAND Corporation, December 1982.
- [11] D. R. Jefferson. Virtual time. *ACM Transactions on Programming Languages and Systems*, 7(3):404--425, July 1985.
- [12] P. L'Ecuyer and T. H. Andres. 'A Random Number Generator Based on the Combination of Four LCGs.' *Mathematics and Computers in Simulation*, volume 44, pages 99--107, 1997.
- [13] R. Ronngren and Rassul Ayani. 'A Comparative Study of Parallel and Sequential Priority Queue Algorithms', *ACM Transactions on Modeling and Computer Simulation*, volume 7, number 2, pages 157--209, April 1997.
- [14] G. R. Yaun, D. W. Bauer, H. L. Bhutada, C. D. Carothers, M. Yuksel, and S. Kalyanaraman, 'Large-scale network simulation techniques: examples of tcp and ospf models,' *SIGCOMM Comput. Commun. Rev.*, vol. 33, no. 3, pp. 27--41, 2003.
- [15] T. Ye and S. Kalyanaraman, 'A unified search framework for large-scale black-box optimization,' Rensselaer Polytechnic Institute, ECSE Department, Networks Lab, Tech. Rep., 2003.
- [16] G. Yaun, C. D. Carothers, and S. Kalyanaraman, 'Large-scale tcp models using optimistic parallel simulation,' in *PADS '03: Proceedings of the seventeenth workshop on Parallel and distributed simulation*. Washington, DC, USA: IEEE Computer Society, 2003, p. 153.